

TK_18_dec

January 5, 2025

1 Data Mining and Machine Learning Lab

1.1 Experiment 1

1.1.1 Importing the necessary libraries

```
[15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

1.1.2 Load the file to your python program in to a data frame DF1.

```
[16]: df1 = pd.read_csv(r"D:\study material\VIT_Data_Science\Winter_Sem\Data Mining_
and Machine Learning Lab\Class_notes\18_Dec_ML_exp1\housepricedata.csv")
```

```
[17]: df1
```

```
[17]:
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
0	8450	7	5	856	2	1	
1	9600	6	8	1262	2	0	
2	11250	7	5	920	2	1	
3	9550	7	5	756	1	0	
4	14260	8	5	1145	2	1	
...	
1455	7917	6	5	953	2	1	
1456	13175	6	6	1542	2	0	
1457	9042	7	9	1152	2	0	
1458	9717	5	6	1078	1	0	
1459	9937	5	6	1256	1	1	

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	GarageArea	AboveMedianPrice
0	3	8	0	548	1
1	3	6	1	460	1
2	3	6	1	608	1
3	3	7	1	642	0
4	4	9	1	836	1
...
1455	3	7	1	460	1

1456	3	7	2	500	1
1457	4	9	2	252	1
1458	2	5	0	240	0
1459	3	6	0	276	0

[1460 rows x 11 columns]

1.1.3 First five observations from your dataset

```
[18]: df1.head(5)
```

```
[18]:   LotArea  OverallQual  OverallCond  TotalBsmtSF  FullBath  HalfBath  \
0    8450             7            5         856         2          1
1    9600             6            8        1262         2          0
2   11250             7            5         920         2          1
3    9550             7            5         756         1          0
4   14260             8            5        1145         2          1

   BedroomAbvGr  TotRmsAbvGrd  Fireplaces  GarageArea  AboveMedianPrice
0              3              8           0          548                1
1              3              6           1          460                1
2              3              6           1          608                1
3              3              7           1          642                0
4              4              9           1          836                1
```

1.1.4 Last five observations from our dataset

```
[19]: df1.tail(5)
```

```
[19]:   LotArea  OverallQual  OverallCond  TotalBsmtSF  FullBath  HalfBath  \
1455    7917             6            5         953         2          1
1456   13175             6            6        1542         2          0
1457    9042             7            9        1152         2          0
1458    9717             5            6        1078         1          0
1459    9937             5            6        1256         1          1

   BedroomAbvGr  TotRmsAbvGrd  Fireplaces  GarageArea  AboveMedianPrice
1455           3              7           1          460                1
1456           3              7           2          500                1
1457           4              9           2          252                1
1458           2              5           0          240                0
1459           3              6           0          276                0
```

1.1.5 Shape of your dataset

```
[20]: df1.shape
```

[20]: (1460, 11)

1.1.6 info of your dataset

[21]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LotArea                1460 non-null   int64
1   OverallQual            1460 non-null   int64
2   OverallCond            1460 non-null   int64
3   TotalBsmtSF            1460 non-null   int64
4   FullBath               1460 non-null   int64
5   HalfBath               1460 non-null   int64
6   BedroomAbvGr           1460 non-null   int64
7   TotRmsAbvGrd           1460 non-null   int64
8   Fireplaces             1460 non-null   int64
9   GarageArea             1460 non-null   int64
10  AboveMedianPrice       1460 non-null   int64
dtypes: int64(11)
memory usage: 125.6 KB
```

1.1.7 Create a new dataframe DF2 with the first 100 observations of your dataset with only the columns LotArea and BedroomAbvGr [Use iloc operator]

[22]: df2 = (df1.iloc[0:100])[['LotArea', 'BedroomAbvGr']]

[23]: df2

```
[23]:   LotArea  BedroomAbvGr
0      8450             3
1      9600             3
2     11250             3
3      9550             3
4     14260             4
..      ...             ...
95     9765             3
96     10264            3
97     10921            3
98     10625             2
99      9320             3
```

[100 rows x 2 columns]

1.1.8 Write or export your dataset DF2 to a csv file DF11.csv and save it.

```
[24]: df2.to_csv('DF11.csv')
```

1.1.9 Find the maximum and minimum of the LotArea column for your dataset DF1

```
[25]: print("Maximum of LotArea column is:",df1['LotArea'].max())
```

Maximum of LotArea column is: 215245

```
[26]: print("Minimum of LotArea column is:",df1['LotArea'].min())
```

Minimum of LotArea column is: 1300

1.1.10 Find the observations from your dataset with LotArea > 10650 from your DF1 dataframe.

```
[27]: df1[df1['LotArea']>10650]
```

```
[27]:
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
2	11250	7	5	920	2	1	
4	14260	8	5	1145	2	1	
5	14115	5	5	796	1	1	
10	11200	5	5	1040	1	0	
11	11924	9	5	1175	3	0	
...		
1442	11003	10	5	1017	2	1	
1446	26142	5	7	1188	1	0	
1448	11767	4	7	560	1	1	
1453	17217	5	5	1140	1	0	
1456	13175	6	6	1542	2	0	
	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	GarageArea	AboveMedianPrice		
2	3	6	1	608		1	
4	4	9	1	836		1	
5	1	5	0	480		0	
10	3	5	0	384		0	
11	4	11	2	736		1	
...		
1442	3	10	1	812		1	
1446	3	6	0	312		0	
1448	2	6	0	384		0	
1453	3	6	0	0		0	
1456	3	7	2	500		1	

[502 rows x 11 columns]

1.1.11 Find the mean, median of your column TotalBsmtSF and find the unique entries (non-repeated ones)

```
[28]: print("Mean of the column TotalBsmtSf is:",df1.TotalBsmtSF.mean())
```

Mean of the column TotalBsmtSf is: 1057.4294520547944

```
[29]: print("Median of the column TotalBsmtSf is:",df1.TotalBsmtSF.median())
```

Median of the column TotalBsmtSf is: 991.5

```
[30]: print("Unique of the column TotalBsmtSf is:\n",df1.TotalBsmtSF.unique())
```

Unique of the column TotalBsmtSf is:

```
[ 856 1262  920  756 1145  796 1686 1107  952  991 1040 1175  912 1494
1253  832 1004    0 1114 1029 1158  637 1777 1060 1566  900 1704 1484
 520  649 1228 1234 1398 1561 1117 1097 1297 1057 1088 1350  840  938
1150 1752 1434 1656  736  955  794  816 1842  384 1425  970  860 1410
 780  530 1370  576 1143 1947 1453  747 1304 2223  845 1086  462  672
1768  440  896 1237 1563 1065 1288  684  612 1013  990 1235  876 1214
 824  680 1588  960  458  950 1610  741 1226 1053  641  789  793 1844
 994 1264 1809 1028  729 1092 1125 1673  728  732 1080 1199 1362 1078
 660 1008  924  992 1063 1267 1461 1907  928  864 1734  910 1490 1728
 715  884  969 1710  825 1602 1200  572  774 1392 1232 1572 1541  882
1149  644 1617 1582  720 1064 1606 1202 1151 1052 2216  968  504 1188
1593  853  725 1431  855 1726 1360  755 1713 1121 1196  617  848 1424
1140 1100 1157 1212  689 1070 1436  686  798 1248 1498 1010  713 2392
 630 1203  483 1373 1194 1462  894 1414  996 1694  735  540  626  948
1845 1020 1367 1444 1573 1302 1314  975 1604  963 1482  506  926 1422
 802  740 1095 1385 1152 1240 1560 2121 1160  807 1468 1575  625  858
 698 1079  768  795 1416 1003  702 1165 1470 2000  700  319  861 1896
 697  972 2136  716 1347 1372 1249 1136 1502 1162  710 1719 1383  844
 596 1056 3206 1358  943 1499 1922 1536 1208 1215  967  721 1684  536
 958 1478  764 1848 1869  616  624  940 1142 1062  888  883 1394 1099
1268  953  744  608  847  683  870 1580 1856  982 1026 1293  939  784
1256  658 1041 1682  804  788 1144  961 1260 1310 1141  806 1281 1034
1276 1340 1344  988  651 1518  907  901  765  799  648 3094 1440 1258
 915 1517  930  813 1533  872 1242 1364  588  709  560 1375 1277 1626
1488  808  547 1976 2153 1705 1833 1792 1216  999 1113 1073  954  264
1269  190 3200  866 1501  777 1218 1368 1084 2006 1244 3138 1379 1257
1452  528 2035  611  707  880 1051 1581 1838 1650  723  654 1204 1069
1709  998  993 1374 1389 1163 1122 1496  846  372 1164 1050 2042 1868
1437  742  770 1722 1814 1430 1058  908  600  965 1032 1299 1120  936
 783 1822 1522  980 1116  978 1156  636 1554 1386  811 1520 1952 1766
 981 1094 2109  525  776 1486 1629 1138 2077 1406 1021 1408  738 1477
2046  923 1291 1195 1190  874  551 1419 2444 1210  927 1112 1391 1800
 360 1473 1643 1324  270  859  718 1176 1311  971 1742  941 1698 1584
1595  868 1153  893 1349 1337 1720 1479 1030 1318 1252  983 1860  836
1935 1614  761 1413  956  712  650  773 1926  731 1417 1024  849 1442
```

```

1649 1568 778 1489 2078 1454 1516 1067 1559 1127 1390 1273 918 1763
1090 1054 1039 1148 1002 1638 105 676 1184 1109 892 2217 1505 1059
951 2330 1670 1623 1017 1105 1001 546 480 1134 1104 1272 1316 1126
1181 1753 964 1466 925 1905 1500 585 1632 819 1616 1161 828 945
979 561 696 1330 817 1098 1428 673 1241 944 1225 1266 1128 485
1930 1396 916 822 750 1700 1007 1187 691 1574 1680 1346 985 1657
602 1022 1082 810 1504 1220 1132 1565 1338 1654 1620 1055 800 1306
1475 2524 1992 1193 973 854 662 1103 1154 942 1048 727 690 1096
1459 1251 1247 1074 1271 290 655 1463 1836 803 833 408 533 1012
1552 1005 1530 974 1567 1006 1042 1298 704 932 1219 1296 1198 959
1261 1598 1683 818 1600 2396 1624 831 1224 663 879 815 1630 2158
931 1660 559 1300 1702 1075 1361 1106 1476 1689 2076 792 2110 1405
1192 746 1986 841 2002 1332 935 1019 661 1309 1328 1085 6110 1246
771 976 1652 1278 1902 1274 1393 1622 1352 420 1795 544 1510 911
693 1284 1732 2033 570 1980 814 873 757 1108 2633 1571 984 1205
714 1746 1525 482 1356 862 839 1286 1485 1594 622 791 708 1223
913 656 1319 1932 539 1221 1542]

```

1.1.12 Sort the dataset DF1 according to the TotalBsmtSF column of your dataset DF1 in ascending and descending order

```

[31]: ##sorting in descending order
df1.sort_values(by=['TotalBsmtSF'],ascending=False)

```

```

[31]:
   LotArea  OverallQual  OverallCond  TotalBsmtSF  FullBath  HalfBath  \
1298    63887           10           5         6110         2         1
332     10655           8           5         3206         2         0
496     12692           8           5         3200         3         0
523     40094          10           5         3138         3         1
440     15431          10           5         3094         2         0
...      ...          ...          ...          ...          ...
1412     7200           4           5           0         2         0
1179     8335           5           5           0         1         0
102      7018           5           5           0         2         0
259     12702           5           5           0         1         0
1048    21750           5           4           0         1         0

   BedroomAbvGr  TotRmsAbvGrd  Fireplaces  GarageArea  AboveMedianPrice
1298           3           12           3         1418           0
332            3            7           1          880           1
496            4           10           1          546           1
523            3           11           1          884           1
440            2           10           2          672           1
...      ...          ...          ...          ...          ...
1412            2            6           0          420           0
1179            3            5           1           0           0
102            4            8           0          410           0

```

259	2	4	0	308	0
1048	3	9	1	336	0

[1460 rows x 11 columns]

```
[32]: #sorted in ascending order
df1.sort_values(by=['TotalBsmtSF'])
```

```
[32]:
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
646	7200	5	5	0	1	0	
1035	11500	4	3	0	1	0	
392	8339	5	7	0	1	0	
749	8405	4	3	0	2	0	
1011	9825	5	5	0	2	0	
...	
440	15431	10	5	3094	2	0	
523	40094	10	5	3138	3	1	
496	12692	8	5	3200	3	0	
332	10655	8	5	3206	2	0	
1298	63887	10	5	6110	2	1	

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	GarageArea	AboveMedianPrice
646	3	7	0	420	0
1035	3	5	0	290	0
392	3	5	0	294	0
749	4	9	0	240	0
1011	4	8	0	0	0
...
440	2	10	2	672	1
523	3	11	1	884	1
496	4	10	1	546	1
332	3	7	1	880	1
1298	3	12	3	1418	0

[1460 rows x 11 columns]

1.1.13 Find the empty cells in 'GarageArea' column of your dataset DF1 and fill it with the average value of the column GarageArea

```
[33]: df1['GarageArea'].isnull().sum()
```

```
[33]: 0
```

```
[34]: df1['GarageArea'].fillna(df1['GarageArea'].mean(),inplace=True)
```

1.1.14 Replace the column named Above median price in your dataframe with 1's where ever you have Yes and 0 where ever you have No


```
[35]: df1['AboveMedianPrice'].replace({'Yes':1, 'no':0}, inplace=True)
```

```
[36]: df1['AboveMedianPrice']
```


```
[36]: 0      1
      1      1
      2      1
      3      0
      4      1
      ..
     1455     1
     1456     1
     1457     1
     1458     0
     1459     0
      Name: AboveMedianPrice, Length: 1460, dtype: int64
```

1.1.15 Draw a scatterplot with columns 'GarageArea' on x axis and 'LotArea' on y-axis

```
[37]: plt.scatter(df1['GarageArea'], df1['LotArea'])
      plt.title("Scatter plot of GarageArea and LotArea")
      plt.xlabel("GarageArea")
      plt.ylabel("LotArea")
      plt.show()
```

```
[38]: import seaborn as sns
sns.scatterplot(x=df1['GarageArea'], y = df1['LotArea'])
plt.title("Scatter plot of GarageArea and LotArea")
plt.xlabel("GarageArea")
plt.ylabel("LotArea")
plt.show()
```



1.1.16 Drop the column 'GarageArea' from your dataset DF1

```
[39]: df1.drop('GarageArea',axis=1,inplace=True)
```

```
[40]: df1
```

```
[40]:
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
0	8450	7	5	856	2	1	
1	9600	6	8	1262	2	0	
2	11250	7	5	920	2	1	
3	9550	7	5	756	1	0	
4	14260	8	5	1145	2	1	
...	
1455	7917	6	5	953	2	1	
1456	13175	6	6	1542	2	0	
1457	9042	7	9	1152	2	0	
1458	9717	5	6	1078	1	0	
1459	9937	5	6	1256	1	1	

BedroomAbvGr	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
--------------	--------------	------------	------------------

0	3	8	0	1
1	3	6	1	1
2	3	6	1	1
3	3	7	1	0
4	4	9	1	1
...
1455	3	7	1	1
1456	3	7	2	1
1457	4	9	2	1
1458	2	5	0	0
1459	3	6	0	0

[1460 rows x 10 columns]

1.2 Q1. Now, normalize the columns of the dataset1 using the above technique and save it to a new csv file DF3.csv

```
[41]: from sklearn import preprocessing
min_max_scaler = preprocessing.MinMaxScaler()
col_name = df1.columns[:]
x = df1.loc[:, col_name]
x = pd.DataFrame(data = min_max_scaler.fit_transform(x), columns = col_name)
print(x)
x.to_csv('df3.csv')
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
0	0.033420	0.666667	0.500	0.140098	0.666667	0.5	
1	0.038795	0.555556	0.875	0.206547	0.666667	0.0	
2	0.046507	0.666667	0.500	0.150573	0.666667	0.5	
3	0.038561	0.666667	0.500	0.123732	0.333333	0.0	
4	0.060576	0.777778	0.500	0.187398	0.666667	0.5	
...	
1455	0.030929	0.555556	0.500	0.155974	0.666667	0.5	
1456	0.055505	0.555556	0.625	0.252373	0.666667	0.0	
1457	0.036187	0.666667	1.000	0.188543	0.666667	0.0	
1458	0.039342	0.444444	0.625	0.176432	0.333333	0.0	
1459	0.040370	0.444444	0.625	0.205565	0.333333	0.5	

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
0	0.375	0.500000	0.000000	1.0
1	0.375	0.333333	0.333333	1.0
2	0.375	0.333333	0.333333	1.0
3	0.375	0.416667	0.333333	0.0
4	0.500	0.583333	0.333333	1.0
...
1455	0.375	0.416667	0.333333	1.0
1456	0.375	0.416667	0.666667	1.0
1457	0.500	0.583333	0.666667	1.0

1458	0.250	0.250000	0.000000	0.0
1459	0.375	0.333333	0.000000	0.0

[1460 rows x 10 columns]

1.2.1 Q2. Now normalize the whole data to the range (2,3) using Min-Max normalization

```
[42]: from sklearn import preprocessing
min_max_scaler = preprocessing.MinMaxScaler(feature_range=(2,3))
col_name = df1.columns[:]
x = df1.loc[:, col_name]
x = pd.DataFrame(data = min_max_scaler.fit_transform(x), columns = col_name)
print(x)
```

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
0	2.033420	2.666667	2.500	2.140098	2.666667	2.5	
1	2.038795	2.555556	2.875	2.206547	2.666667	2.0	
2	2.046507	2.666667	2.500	2.150573	2.666667	2.5	
3	2.038561	2.666667	2.500	2.123732	2.333333	2.0	
4	2.060576	2.777778	2.500	2.187398	2.666667	2.5	
...	
1455	2.030929	2.555556	2.500	2.155974	2.666667	2.5	
1456	2.055505	2.555556	2.625	2.252373	2.666667	2.0	
1457	2.036187	2.666667	3.000	2.188543	2.666667	2.0	
1458	2.039342	2.444444	2.625	2.176432	2.333333	2.0	
1459	2.040370	2.444444	2.625	2.205565	2.333333	2.5	

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
0	2.375	2.500000	2.000000	3.0
1	2.375	2.333333	2.333333	3.0
2	2.375	2.333333	2.333333	3.0
3	2.375	2.416667	2.333333	2.0
4	2.500	2.583333	2.333333	3.0
...
1455	2.375	2.416667	2.333333	3.0
1456	2.375	2.416667	2.666667	3.0
1457	2.500	2.583333	2.666667	3.0
1458	2.250	2.250000	2.000000	2.0
1459	2.375	2.333333	2.000000	2.0

[1460 rows x 10 columns]

1.3 Q3. Now do decimal scaling for the original column data of the column LotArea of your initial dataframe and print the results.

```
[43]: max_abs_values = abs(df1.max())
```

```
[44]: max_abs_values
```

```
[44]: LotArea          215245
      OverallQual      10
      OverallCond       9
      TotalBsmtSF      6110
      FullBath          3
      HalfBath          2
      BedroomAbvGr       8
      TotRmsAbvGrd      14
      Fireplaces         3
      AboveMedianPrice    1
      dtype: int64
```

```
[45]: df88 = df1['LotArea'].apply(lambda x : x/
      ↪10**(len(str(max_abs_values['LotArea']))))
```

```
[46]: df88
```

```
[46]: 0      0.008450
      1      0.009600
      2      0.011250
      3      0.009550
      4      0.014260
      ...
      1455    0.007917
      1456    0.013175
      1457    0.009042
      1458    0.009717
      1459    0.009937
      Name: LotArea, Length: 1460, dtype: float64
```

1.3.1 Q4. Now try to standardize the whole data in the dataframe and print the dataframe.

```
[47]: from sklearn import preprocessing
      standard_scaler = preprocessing.StandardScaler()
      col_name = df1.columns
      x = df1.loc[:, col_name]
      x = pd.DataFrame(data = standard_scaler.fit_transform(x), columns = col_name)
      print(x)
```

```
      LotArea  OverallQual  OverallCond  TotalBsmtSF  FullBath  HalfBath  \
0    -0.207142    0.651479   -0.517200   -0.459303    0.789741    1.227585
```

1	-0.091886	-0.071836	2.179628	0.466465	0.789741	-0.761621
2	0.073480	0.651479	-0.517200	-0.313369	0.789741	1.227585
3	-0.096897	0.651479	-0.517200	-0.687324	-1.026041	-0.761621
4	0.375148	1.374795	-0.517200	0.199680	0.789741	1.227585
...
1455	-0.260560	-0.071836	-0.517200	-0.238122	0.789741	1.227585
1456	0.266407	-0.071836	0.381743	1.104925	0.789741	-0.761621
1457	-0.147810	0.651479	3.078570	0.215641	0.789741	-0.761621
1458	-0.080160	-0.795151	0.381743	0.046905	-1.026041	-0.761621
1459	-0.058112	-0.795151	0.381743	0.452784	-1.026041	1.227585

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
0	0.163779	0.912210	-0.951226	1.002743
1	0.163779	-0.318683	0.600495	1.002743
2	0.163779	-0.318683	0.600495	1.002743
3	0.163779	0.296763	0.600495	-0.997264
4	1.390023	1.527656	0.600495	1.002743
...
1455	0.163779	0.296763	0.600495	1.002743
1456	0.163779	0.296763	2.152216	1.002743
1457	1.390023	1.527656	2.152216	1.002743
1458	-1.062465	-0.934130	-0.951226	-0.997264
1459	0.163779	-0.318683	-0.951226	-0.997264

[1460 rows x 10 columns]

[48]: x

[48]:

	LotArea	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	\
0	-0.207142	0.651479	-0.517200	-0.459303	0.789741	1.227585	
1	-0.091886	-0.071836	2.179628	0.466465	0.789741	-0.761621	
2	0.073480	0.651479	-0.517200	-0.313369	0.789741	1.227585	
3	-0.096897	0.651479	-0.517200	-0.687324	-1.026041	-0.761621	
4	0.375148	1.374795	-0.517200	0.199680	0.789741	1.227585	
...	
1455	-0.260560	-0.071836	-0.517200	-0.238122	0.789741	1.227585	
1456	0.266407	-0.071836	0.381743	1.104925	0.789741	-0.761621	
1457	-0.147810	0.651479	3.078570	0.215641	0.789741	-0.761621	
1458	-0.080160	-0.795151	0.381743	0.046905	-1.026041	-0.761621	
1459	-0.058112	-0.795151	0.381743	0.452784	-1.026041	1.227585	

	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
0	0.163779	0.912210	-0.951226	1.002743
1	0.163779	-0.318683	0.600495	1.002743
2	0.163779	-0.318683	0.600495	1.002743
3	0.163779	0.296763	0.600495	-0.997264
4	1.390023	1.527656	0.600495	1.002743

...
1455	0.163779	0.296763	0.600495	1.002743
1456	0.163779	0.296763	2.152216	1.002743
1457	1.390023	1.527656	2.152216	1.002743
1458	-1.062465	-0.934130	-0.951226	-0.997264
1459	0.163779	-0.318683	-0.951226	-0.997264

[1460 rows x 10 columns]

1.4 Train Test splitting of data for model training.

1.4.1 Now perform 70:30 train test split for our dataframe data with the target variable or output variable as LotArea and print the training and testing data which may be we can use for fitting a model for this data. Say like trying to predict the LotArea for a house based on all the other features.

```
[49]: from sklearn.model_selection import train_test_split
x = df1.drop('LotArea', axis=1)
y = df1['LotArea']
#Split the data into training and testing sets (70% train, 30% test)
x_train,x_test,y_train,y_test = train_test_split(x, y, test_size=0.3,
↳random_state=42)
#Print the testing and training data
print("Training Features:\n", x_train)
print("Testing Features:\n", x_test)
print("Training Target:\n", y_train)
print("Testing Target:\n", y_test)
```

Training Features:

	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	BedroomAbvGr
\						
135	7	6	1304	2	0	3
1452	5	5	547	1	0	2
762	7	5	756	2	1	3
932	9	5	1905	2	0	3
435	7	6	799	2	1	3
...
1095	6	5	1314	2	0	3
1130	4	3	1122	2	0	4
1294	5	7	864	1	0	2
860	7	8	912	1	1	3
1126	7	5	1373	2	0	2

	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
135	7	1	1
1452	5	0	0
762	7	0	1
932	8	1	1

435	6	1	1
...
1095	6	1	1
1130	7	2	0
1294	5	0	0
860	7	1	1
1126	7	1	1

[1022 rows x 9 columns]

Testing Features:

	OverallQual	OverallCond	TotalBsmtSF	FullBath	HalfBath	BedroomAbvGr
\						
892	6	8	1059	1	0	3
1105	8	5	1463	2	1	3
413	5	6	1008	1	0	2
522	6	7	1004	2	0	3
1036	9	5	1620	2	0	2
...
331	5	6	1056	1	0	3
323	3	8	1162	1	0	3
650	7	6	813	2	1	3
439	6	8	684	1	0	3
798	9	5	1926	3	1	4

	TotRmsAbvGrd	Fireplaces	AboveMedianPrice
892	6	0	0
1105	9	2	1
413	5	1	0
522	7	2	0
1036	6	1	1
...
331	6	0	0
323	6	0	0
650	7	0	1
439	7	0	0
798	11	2	1

[438 rows x 9 columns]

Training Target:

135	10400
1452	3675
762	8640
932	11670
435	10667
...	...
1095	9317
1130	7804
1294	8172


```

860      7642
1126      3684
Name: LotArea, Length: 1022, dtype: int64
Testing Target:
  892      8414
1105     12256
  413      8960
  522      5000
1036     12898
...
 331      8176
 323      5820
 650      8125
 439     12354
 798     13518
Name: LotArea, Length: 438, dtype: int64

```

```

[50]: # Linear Regression model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
model = LinearRegression()
model.fit(x_train, y_train)

# Predicting the LotArea for test data
y_pred = model.predict(x_test)

# Evaluating the model's performance
mse = mean_squared_error(y_test, y_pred) # Mean Squared Error
r2 = r2_score(y_test, y_pred)           # R-squared score

print("Model Coefficients:", model.coef_)
print("Model Intercept:", model.intercept_)
print("Mean Squared Error:", mse)
print("R-squared Score:", r2)

results = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
print(results)

```

```

Model Coefficients: [-1784.57317854  535.39223068    6.34890747  105.64506875
 -350.6656404   499.86480782  514.75964975  3611.59648719
 2444.37464586]

```

```

Model Intercept: 3657.748961708652

```

```

Mean Squared Error: 31147397.821340438

```

```

R-squared Score: 0.03359678635359464

```

	Actual	Predicted
892	8414	8650.738138
1105	12256	17007.199705
413	8960	11637.704604

522	5000	15609.753689
1036	12898	10914.330395
...
331	8176	9345.480133
323	5820	14658.395143
650	8125	6947.662984
439	12354	6784.657486
798	13518	19797.199861

[438 rows x 2 columns]