

# Census\_income\_prediction

April 2, 2025

## 1 Data Mining and Machine Learning (PMDS505L)

### 1.1 Digital Assignment-1

## 2 Problem Statement

The objective of this project is to apply multiple binary classification algorithms to a selected dataset to predict a target variable with two possible outcomes.

For this project, I have chosen a [Census Income Dataset](#) where the goal is to predict whether the annual income of an individual exceeds \$50K/yr based on census data

### 2.1 Dataset Information

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
age	Feature	Integer	Age	N/A		no
workclass	Feature	Categorical	Income	Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.		yes
fnlwgt	Feature	Integer				no

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
education	Feature	Categorical	Education Level	Bachelors, Some- college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.		no
education- num	Feature	Integer	Education Level			no
marital-status	Feature	Categorical	Other	Married-civ- spouse, Divorced, Never- married, Separated, Widowed, Married- spouse- absent, Married-AF- spouse.		no

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
occupation	Feature	Categorical	Other	Tech- support, Craft-repair, Other- service, Sales, Exec- managerial, Prof- specialty, Handlers- cleaners, Machine-op- inspct, Adm- clerical, Farming- fishing, Transport- moving, Priv-house- serv, Protective- serv, Armed- Forces.		yes
relationship	Feature	Categorical	Other	Wife, Own-child, Husband, Not-in- family, Other- relative, Unmarried.		no
race	Feature	Categorical	Race	White, Asian-Pac- Islander, Amer- Indian- Eskimo, Other, Black.		no
sex	Feature	Binary	Sex	Female, Male.		no
capital-gain	Feature	Integer				no
capital-loss	Feature	Integer				no

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
hours-per-week	Feature	Integer				no

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
native-country	Feature	Categorical	Other	United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Tri-nadad&Tobago, Peru, Hong, Holand-Netherlands.		yes

Variable Name	Role	Type	Demographic	Description	Units	Missing Values
income	Target	Binary	Income	>50K, <=50K.		no

### 2.1.1 Loading the necessary libraries

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, \
    GradientBoostingClassifier, BaggingClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score, \
    recall_score, classification_report, roc_auc_score, \
    confusion_matrix, roc_curve, auc
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE

import warnings
warnings.filterwarnings("ignore")
```

### 2.1.2 Loading the dataset

```
[2]: url = "https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.
    data"
colnames = ['age', 'workclass', 'fnlwgt', 'education', 'education-num', \
    'marital-status', 'occupation', \
    'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', \
    'hours-per-week', 'native-country', 'income']
data = pd.read_csv(url, header=None, names=colnames, na_values=' ?')
data.head()
```

```
[2]:   age      workclass  fnlwgt  education  education-num \
0   39      State-gov   77516   Bachelors             13
1   50  Self-emp-not-inc  83311   Bachelors             13
2   38      Private   215646    HS-grad              9
3   53      Private   234721      11th              7
4   28      Private   338409   Bachelors             13
```

	marital-status	occupation	relationship	race	sex \
0	Never-married	Adm-clerical	Not-in-family	White	Male
1	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[3]: print("Shape of the dataset is:", data.shape)
```

Shape of the dataset is: (32561, 15)

### 2.1.3 Checking for missing values

```
[4]: data.isnull().sum()
```

```
[4]: age                0
workclass             1836
fnlwgt                0
education              0
education-num         0
marital-status        0
occupation            1843
relationship          0
race                  0
sex                   0
capital-gain          0
capital-loss          0
hours-per-week        0
native-country        583
income                0
dtype: int64
```

### 2.1.4 Dropping the missing values

```
[5]: data.dropna(inplace = True)
```

```
[6]: data.isna().sum().sum()
```

```
[6]: 0
```

### 2.1.5 Checking for duplicate values

```
[7]: data.duplicated().sum()
```

```
[7]: 23
```

### 2.1.6 Dropping the duplicate values

```
[8]: data.drop_duplicates(inplace = True)
```

```
[9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 30139 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30139 non-null  int64
1   workclass              30139 non-null  object
2   fnlwgt                 30139 non-null  int64
3   education              30139 non-null  object
4   education-num          30139 non-null  int64
5   marital-status         30139 non-null  object
6   occupation             30139 non-null  object
7   relationship           30139 non-null  object
8   race                   30139 non-null  object
9   sex                   30139 non-null  object
10  capital-gain           30139 non-null  int64
11  capital-loss           30139 non-null  int64
12  hours-per-week         30139 non-null  int64
13  native-country         30139 non-null  object
14  income                 30139 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

### 2.1.7 Setting the feature and target variables

```
[10]: x = data.drop('income', axis=1)
      y = (data['income'] == '>50K').astype(int) # 1 = >50K, 0 = <=50K
```

```
[11]: y.value_counts(normalize=True)
```

```
[11]: income
0    0.750954
1    0.249046
Name: proportion, dtype: float64
```

- Data is imbalanced



### 2.1.8 Encoding the Categorical Variables

```
[12]: categorical_cols = ['workclass', 'education', 'marital-status', 'occupation',
    ↪ 'relationship', 'race', 'sex', 'native-country']
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    x[col] = le.fit_transform(x[col])
    label_encoders[col] = le
```

```
[13]: x
```

```
[13]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	39	5	77516	9	13	4	
1	50	4	83311	9	13	2	
2	38	2	215646	11	9	0	
3	53	2	234721	1	7	2	
4	28	2	338409	9	13	2	
...	...	...	...	...	...	...	
32556	27	2	257302	7	12	2	
32557	40	2	154374	11	9	2	
32558	58	2	151910	11	9	6	
32559	22	2	201490	11	9	4	
32560	52	3	287927	11	9	2	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	0	1	4	1	2174	0	
1	3	0	4	1	0	0	
2	5	1	4	1	0	0	
3	5	0	2	1	0	0	
4	9	5	2	0	0	0	
...	...	...	...	...	...	...	
32556	12	5	4	0	0	0	
32557	6	0	4	1	0	0	
32558	0	4	4	0	0	0	
32559	0	3	4	1	0	0	
32560	3	5	4	0	15024	0	

	hours-per-week	native-country
0	40	38
1	13	38
2	40	38
3	40	38
4	40	4
...	...	...
32556	38	38
32557	40	38

32558	40	38
32559	20	38
32560	40	38

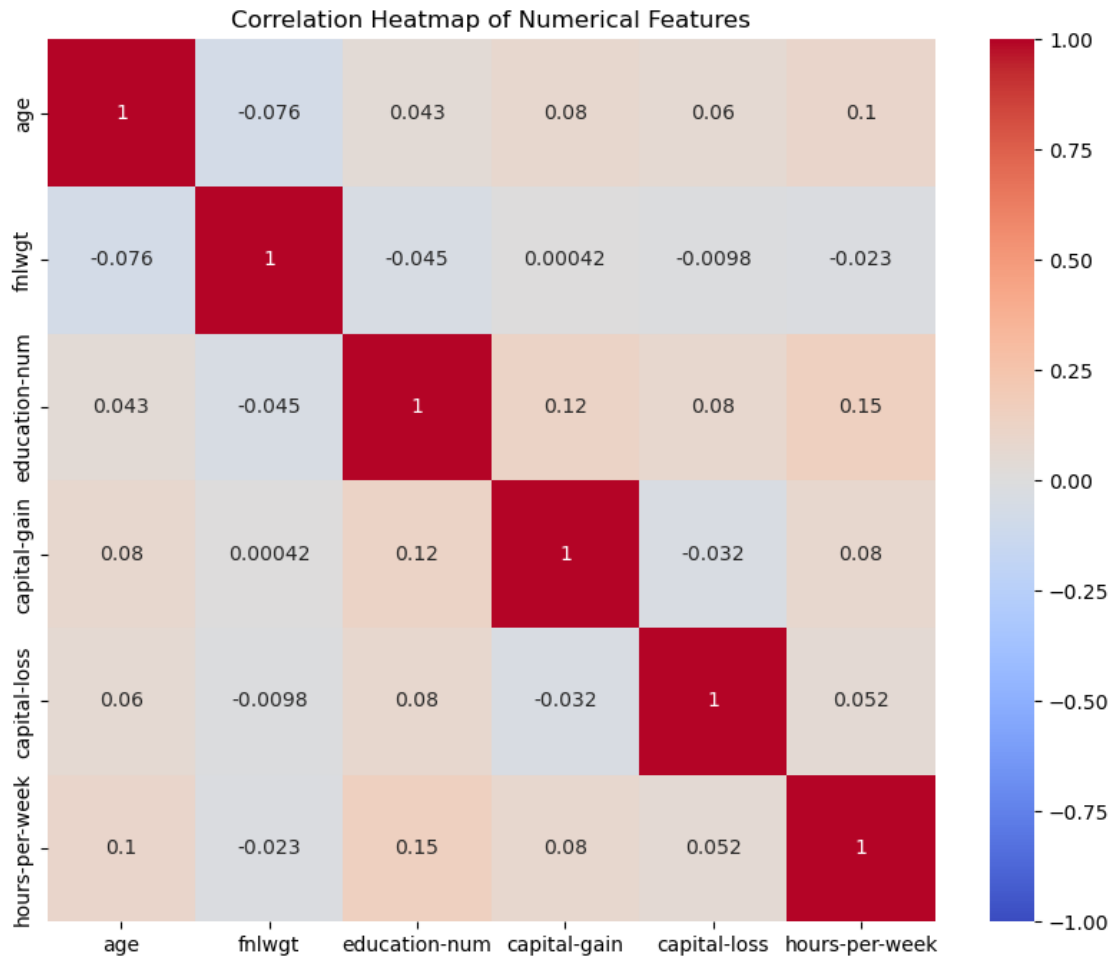
[30139 rows x 14 columns]

### 2.1.9 Scaling Numerical Features

```
[14]: numerical_cols = ['age', 'fnlwgt', 'education-num', 'capital-gain',  
    ↪ 'capital-loss', 'hours-per-week']  
scaler = StandardScaler()  
x[numerical_cols] = scaler.fit_transform(x[numerical_cols])
```

### 2.1.10 Correlation Heatmap for numerical features

```
[15]: plt.figure(figsize=(10, 8))  
sns.heatmap(data[numerical_cols].corr(), annot=True, cmap='coolwarm', vmin=-1,  
    ↪ vmax=1, center=0)  
plt.title('Correlation Heatmap of Numerical Features')  
plt.show()
```



- Very high correlation is not present among the variables

### 2.1.11 Train test split of the data

```
[16]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,
↳random_state=42,stratify=y)
```

### 2.1.12 Applying SMOTE for class balancing

```
[17]: smote = SMOTE(random_state=42)
x_train_smote, y_train_smote = smote.fit_resample(x_train, y_train)

print("\nClass distribution after SMOTE:")
print(y_train_smote.value_counts(normalize=True))
```

Class distribution after SMOTE:  
income

```
0    0.5
1    0.5
Name: proportion, dtype: float64
```

### 2.1.13 Model building Function

```
[18]: def train_and_evaluate_models(x_train, y_train, x_test, y_test):
    # Defining the classification models
    classifiers = {
        'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
        'Decision Tree': DecisionTreeClassifier(random_state=42),
        'Random Forest': RandomForestClassifier(random_state=42),
        'AdaBoost': AdaBoostClassifier(random_state=42),
        'Gradient Boosting': GradientBoostingClassifier(random_state=42),
        'Bagging': BaggingClassifier(random_state=42),
        'SVC': SVC(probability=True, random_state=42, kernel='linear')
    }

    # Storing the results results
    results = {}

    for name, clf in classifiers.items():
        clf.fit(x_train, y_train)
        y_pred = clf.predict(x_test)
        results[name] = classification_report(y_test, y_pred, output_dict=True)
        print(f"\n{name}:\n", classification_report(y_test, y_pred))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(f'Confusion Matrix - {name}')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()

    # ROC Curve
    y_prob = clf.predict_proba(x_test)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    plt.figure(figsize=(6, 4))
    plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve - {name}')
    plt.legend(loc='best')
    plt.show()
```

```


    return results


# Calling the function
results = train_and_evaluate_models(x_train_smote, y_train_smote, x_test,
    ↪y_test)

```

Logistic Regression:



	precision	recall	f1-score	support
0	0.90	0.77	0.83	4527
1	0.52	0.74	0.61	1501
accuracy			0.77	6028
macro avg	0.71	0.76	0.72	6028
weighted avg	0.81	0.77	0.78	6028






Decision Tree:

	precision	recall	f1-score	support
0	0.87	0.84	0.86	4527
1	0.57	0.63	0.60	1501
accuracy			0.79	6028
macro avg	0.72	0.74	0.73	6028
weighted avg	0.80	0.79	0.79	6028




Random Forest:

	precision	recall	f1-score	support
0	0.90	0.88	0.89	4527
1	0.66	0.72	0.69	1501
accuracy			0.84	6028
macro avg	0.78	0.80	0.79	6028
weighted avg	0.84	0.84	0.84	6028










AdaBoost:


	precision	recall	f1-score	support
0	0.92	0.83	0.87	4527
1	0.60	0.79	0.69	1501
accuracy			0.82	6028
macro avg	0.76	0.81	0.78	6028
weighted avg	0.84	0.82	0.83	6028



Gradient Boosting:



	precision	recall	f1-score	support
0	0.93	0.83	0.88	4527
1	0.61	0.80	0.70	1501
accuracy			0.82	6028
macro avg	0.77	0.82	0.79	6028
weighted avg	0.85	0.82	0.83	6028






Bagging:


	precision	recall	f1-score	support
0	0.89	0.89	0.89	4527
1	0.66	0.65	0.66	1501
accuracy			0.83	6028
macro avg	0.77	0.77	0.77	6028
weighted avg	0.83	0.83	0.83	6028



SVC:

	precision	recall	f1-score	support
0	0.91	0.76	0.83	4527
1	0.51	0.76	0.61	1501
accuracy			0.76	6028
macro avg	0.71	0.76	0.72	6028
weighted avg	0.81	0.76	0.77	6028





#### 2.1.14 Model Comparison

```
[19]: top_models = sorted(results.items(), key=lambda x: x[1]['1']['f1-score'],
    ↪reverse=True)[:3] # Top 3 models
    print("\nTop Models for Tuning:", [name for name, _ in top_models])
```

Top Models for Tuning: ['Gradient Boosting', 'Random Forest', 'AdaBoost']

#### 2.1.15 Hyper parameter Tuning for the Top performing Models

```
[20]: classifiers = {
    'Random Forest': RandomForestClassifier(random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),
    'AdaBoost': AdaBoostClassifier(random_state=42)
}

param_grids = {
    'Random Forest':
        {
            'n_estimators': [100, 200, 400, 500],
            'max_depth': [10, 20, None],
            'min_samples_split': [2, 5],
```

```

        'min_samples_leaf': [1, 2]
    },
    'Gradient Boosting':
    {
        'n_estimators': [100, 200, 300, 500, 700],
        'learning_rate': [0.01, 0.05, 0.1, 0.5],
        'max_depth': [3, 5]
    },
    'AdaBoost':
    {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.05, 0.1, 0.5, 1.0]
    }
}

tuned_models = {}
for name in ['Random Forest', 'Gradient Boosting', 'AdaBoost']:
    print(f"\nTuning {name}...")
    grid = GridSearchCV(
        classifiers[name],
        param_grids[name],
        cv=5,
        scoring='f1',
        n_jobs=-1
    )
    grid.fit(x_train_smote, y_train_smote)
    tuned_models[name] = grid.best_estimator_

    # Report best parameters and performance
    print(f"Tuned {name} Best Params: {grid.best_params_}")
    y_pred_tuned = grid.predict(x_test)
    print(f"Tuned {name} Performance:\n", classification_report(y_test,
↪y_pred_tuned))
    # Select and save the best model
    best_model_name = max(tuned_models, key=lambda x: classification_report(y_test,
↪tuned_models[x].predict(x_test), output_dict=True)['1']['f1-score'])
    best_model = tuned_models[best_model_name]
    y_pred_final = best_model.predict(x_test)
    print(f"\nBest Model ({best_model_name}) Final Performance:\n",
↪classification_report(y_test, y_pred_final))

```

Tuning Random Forest...

Tuned Random Forest Best Params: {'max\_depth': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 500}

Tuned Random Forest Performance:

precision	recall	f1-score	support
-----------	--------	----------	---------



0	0.90	0.88	0.89	4527
1	0.67	0.72	0.69	1501
accuracy			0.84	6028
macro avg	0.79	0.80	0.79	6028
weighted avg	0.84	0.84	0.84	6028

Tuning Gradient Boosting...

Tuned Gradient Boosting Best Params: {'learning\_rate': 0.1, 'max\_depth': 5, 'n\_estimators': 300}

Tuned Gradient Boosting Performance:

	precision	recall	f1-score	support
0	0.91	0.90	0.90	4527
1	0.71	0.73	0.72	1501
accuracy			0.86	6028
macro avg	0.81	0.81	0.81	6028
weighted avg	0.86	0.86	0.86	6028

Tuning AdaBoost...

Tuned AdaBoost Best Params: {'learning\_rate': 1.0, 'n\_estimators': 200}

Tuned AdaBoost Performance:

	precision	recall	f1-score	support
0	0.92	0.84	0.88	4527
1	0.62	0.79	0.70	1501
accuracy			0.83	6028
macro avg	0.77	0.82	0.79	6028
weighted avg	0.85	0.83	0.83	6028

Best Model (Gradient Boosting) Final Performance:

	precision	recall	f1-score	support
0	0.91	0.90	0.90	4527
1	0.71	0.73	0.72	1501
accuracy			0.86	6028
macro avg	0.81	0.81	0.81	6028
weighted avg	0.86	0.86	0.86	6028


### 2.1.16 Plotting decision boundaries

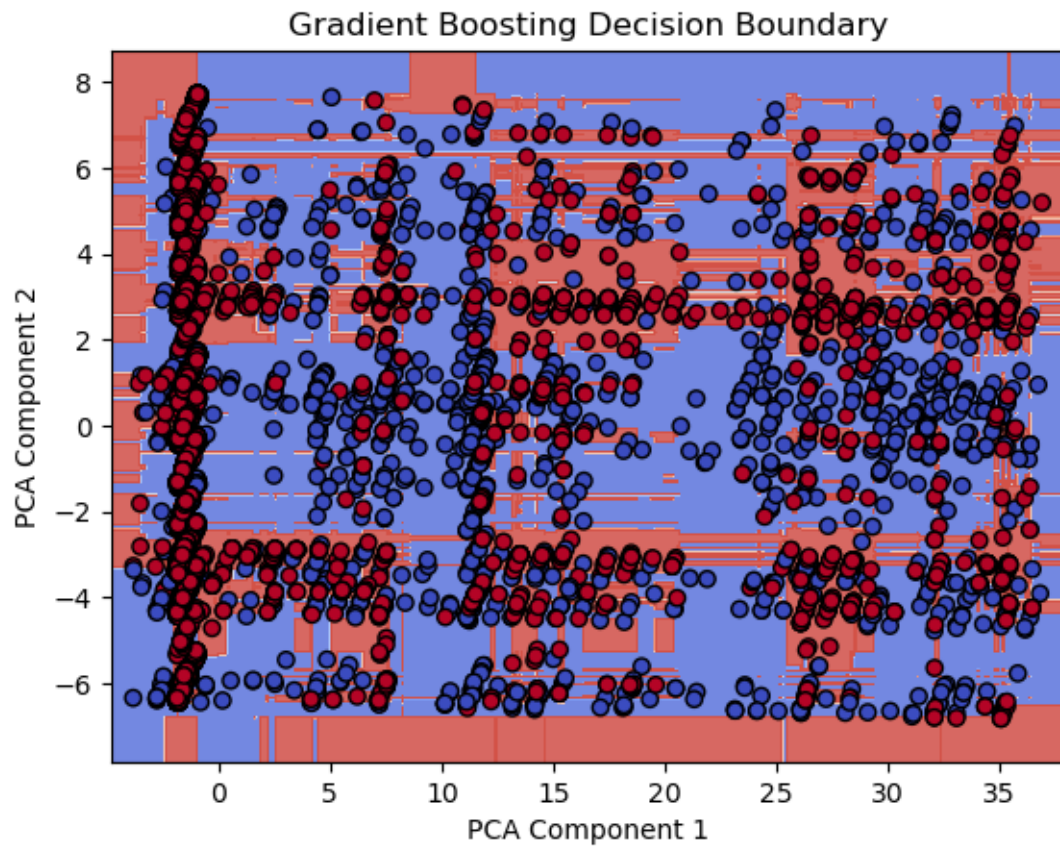
#### Using PCA to reduce data to 2D

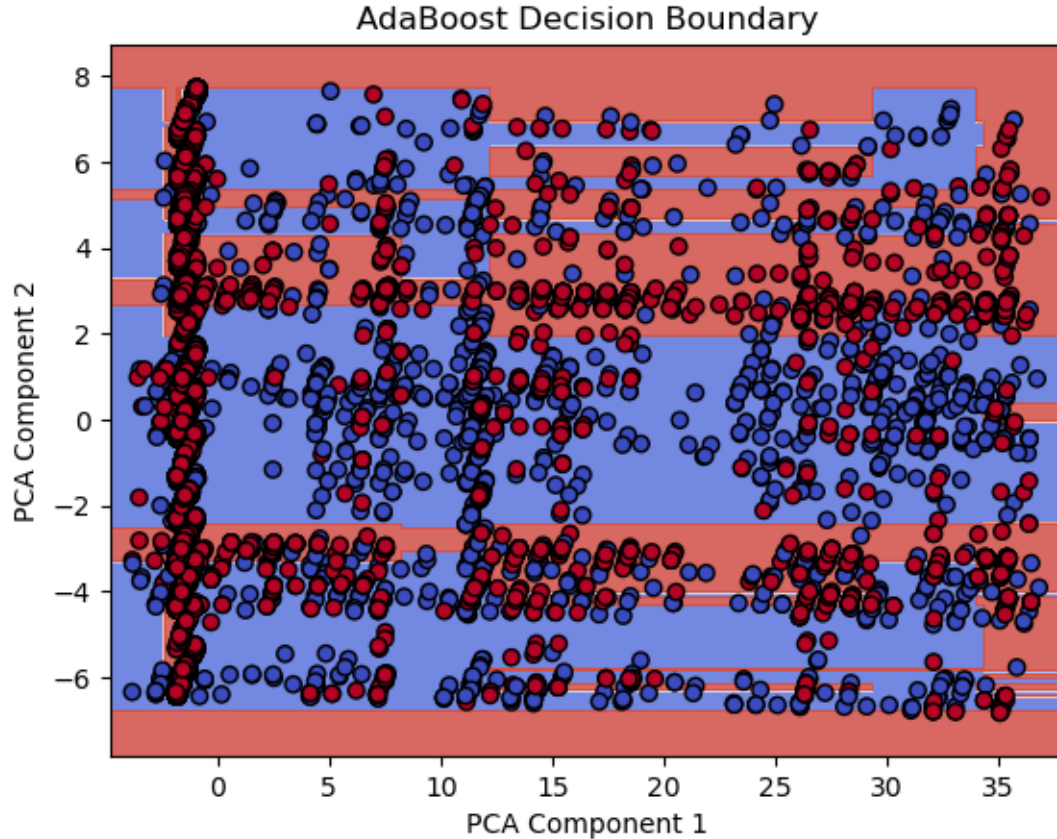
```
[21]: pca = PCA(n_components=2)
x_train_2d = pca.fit_transform(x_train_smote)
x_test_2d = pca.transform(x_test)

def plot_decision_boundary(clf, X, y, title):
    clf.fit(X, y)
    h = 0.02
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, cmap='coolwarm', alpha=0.8)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap='coolwarm', edgecolors='k')
    plt.title(title)
    plt.xlabel('PCA Component 1')
    plt.ylabel('PCA Component 2')
    plt.show()

# Plot decision boundaries for tuned models
for name, model in tuned_models.items():
    plot_decision_boundary(model, x_train_2d, y_train_smote, f'{name} Decision_
↳Boundary')
```







## 2.2 Conclusion:

- Model Exploration: Seven classifiers were initially evaluated— **Logistic Regression** (accuracy: 0.77), **Decision Tree** (0.79), **Random Forest** (0.84), **AdaBoost** (0.82), **Gradient Boosting** (0.82), **Bagging** (0.83), and **SVC** (0.76) with Random Forest and Gradient Boosting showing the highest F1-scores for the minority class (>50K).
- Hyperparameter Tuning: The top three models (Gradient Boosting, Random Forest, AdaBoost) were tuned using GridSearchCV with F1-score as the metric:
- Tuned Random Forest: Achieved an accuracy of 0.84 and F1-score of 0.69 for >50K (params: `n_estimators=500`, `max_depth=None`, `min_samples_split=2`, `min_samples_leaf=1`).
- Tuned Gradient Boosting: Outperformed others with an accuracy of 0.86 and F1-score of 0.72 for >50K (params: `n_estimators=300`, `learning_rate=0.1`, `max_depth=5`).
- Tuned AdaBoost: Recorded an accuracy of 0.83 and F1-score of 0.70 for >50K (params: `n_estimators=200`, `learning_rate=1.0`).
- Best Model: Gradient Boosting was selected as the best model, offering the highest F1-score (0.72) and accuracy (0.86), with strong precision (0.71) and recall (0.73) for identifying high-income individuals (>50K).

The tuned Gradient Boosting model effectively predicts whether an individual's income exceeds \$50K/year, making it suitable for applications like socioeconomic analysis or targeted resource allocation.