

Experiment_7_19_Feb

February 19, 2025

1 Data Mining and machine Learning

2 Experiment 7

2.1 19 February

3 Name: Tufan Kundu

4 Reg no.: 24MDT0184

5 Decision Tree: Bagging, Boosting, RandomForestClassifier

5.1 Q1. Today we will try to see how bagging, boosting etc can be implemented.

```
[86]: ## Loading the necessary libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
```

```
[87]: ## Loading the dataset

df = pd.read_csv(r"C:\Users\Batch1\Documents\Downloads\TK\12_feb\liver_patient.
↪csv")
df
```

```
[87]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	\
0	65	Female	0.7	0.1	187	
1	62	Male	10.9	5.5	699	
2	62	Male	7.3	4.1	490	
3	58	Male	1.0	0.4	182	
4	72	Male	3.9	2.0	195	
..	
578	60	Male	0.5	0.1	500	
579	40	Male	0.6	0.1	98	
580	52	Male	0.8	0.2	245	
581	31	Male	1.3	0.5	184	

582	38	Male	1.0	0.3	216
-----	----	------	-----	-----	-----

	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	\
0	16	18	6.8	
1	64	100	7.5	
2	60	68	7.0	
3	14	20	6.8	
4	27	59	7.3	
..	
578	20	34	5.9	
579	35	31	6.0	
580	48	49	6.4	
581	29	32	6.8	
582	21	24	7.3	

	Albumin	Albumin_and_Globulin_Ratio	liver_disease
0	3.3	0.90	1
1	3.2	0.74	1
2	3.3	0.89	1
3	3.4	1.00	1
4	2.4	0.40	1
..
578	1.6	0.37	0
579	3.2	1.10	1
580	3.2	1.00	1
581	3.4	1.00	1
582	4.4	1.50	0

[583 rows x 11 columns]

5.1.1 Dropping the unnecessary Age and Gender column

```
[88]: df.drop(['Age', 'Gender'], axis=1, inplace = True)
```

5.1.2 Perform Min-Max scaling

```
[89]: from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
X = scaler.fit_transform(df)
```

```
[90]: x = X[:, :-1]
y = X[:, -1]
```

5.2 Do the train test split of the data with test size 20%.

```
[91]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=.
      ↪2,random_state=0)
```

5.2.1 Fit the LogisticRegression model to the this training data.

```
[92]: model_lr = LogisticRegression()
      model_lr.fit(x_train,y_train)
      y_pred_lr = model.predict(x_test)
      print(f"The accuracy of the logistic regression model is:
      ↪{round(accuracy_score(y_test,y_pred_lr)*100,3)} %")
```

The accuracy of the logistic regression model is:66.667 %

5.3 Next for the same data fit a Decision Tree with the same training data and check for the testing accuracy.

```
[93]: from sklearn.tree import DecisionTreeClassifier

      model_dc = DecisionTreeClassifier()
      model_dc.fit(x_train,y_train)
      y_pred_dc = model_dc.predict(x_test)
      print(f"Accuracy score of decision tree:
      ↪{round(accuracy_score(y_test,y_pred_dc)*100,2)}%")
```

Accuracy score of decision tree:60.68%

5.3.1 - The accuracy of the logistic regression model is:66.667 % %

5.3.2 - Accuracy score of decision tree:60.68%

6 Bagging: Bagging Classifier and Bagging regressor.

```
[94]: from sklearn.ensemble import BaggingClassifier, RandomForestClassifier
      BC = BaggingClassifier(n_estimators = 100,random_state = 0)
      #n_estimator is number of trees we are fitting for bagging
      BC.fit(x_train,y_train)
      print(BC.predict(x_test))
      print(accuracy_score(BC.predict(x_test),y_test)*100)
```

```
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1.
 1. 0. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0. 0. 1. 1. 1.
 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 0. 1.
 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1.
 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 0. 1. 0. 1. 0. 1. 1. 1.]
```

67.52136752136752

6.1 For pasting

```
[95]: PBC = BaggingClassifier(n_estimators = 100, random_state = 0, bootstrap=False)
      #bootstrap = False is samples are drawn with replacement
      PBC.fit(x_train, y_train)
      print(PBC.predict(x_test))
      print(accuracy_score(PBC.predict(x_test), y_test)*100)
```

```
[0. 1. 1. 1. 1. 1. 1. 0. 1. 0. 0. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0.
 1. 0. 1. 1. 1. 0. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0. 0. 1. 1. 1. 1.
 1. 1. 1. 0. 1. 0. 1. 0. 1. 0. 1. 0. 1. 1. 1. 0. 1. 1. 1. 1. 0. 1. 0. 1.
 1. 1. 1. 1. 1. 1. 0. 0. 1. 1. 1. 1. 0. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1.
 1. 0. 1. 1. 0. 1. 1. 0. 1. 1. 1. 1. 0. 0. 0. 0. 0. 0. 1. 0. 1.]
62.39316239316239
```

6.2 Implementation of Random Forest classifier

```
[96]: RFC = RandomForestClassifier(n_estimators=100,
      ↪max_features='sqrt', random_state=0)
      RFC.fit(x_train, y_train)
      print(RFC.predict(x_test))
      print(accuracy_score(RFC.predict(x_test), y_test)*100)
```

```
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 0. 0. 1. 1. 1. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 0. 1.
 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 0. 1.
 1. 0. 1. 1. 1. 1. 0. 1. 1. 1. 1. 1. 1. 1. 1. 1. 0. 1. 0. 1.]
68.37606837606837
```

6.3 Checking the crucial features

```
[97]: print(RFC.feature_importances_)
```

```
[0.11831252 0.07608059 0.17098602 0.16769726 0.1466726  0.11859783
 0.1066352  0.09501798]
```

6.4 - Features with large score were crucial in modelling

6.5 Q2. Form a synthetic dataset using the make_classification class which we have used in the previous labs with two features and 3 classes. Fit the decisiontreeclassifier, baggingclassifier and randomforestclassifier and print the decision boundaries for the different classifiers. Visualizing in the previous case is not possible so we can see how the nonlinear boundaries are getting created in these models except for a linear boundary that we have seen by other classifiers in the previous labs.

```
[98]: from sklearn.datasets import make_classification
      # generating a fake data first
      x, y = make_classification(
          n_samples=200, # total number of samples in the dataset
          n_features=2,
          n_classes= 3,
          n_redundant=0,
          n_clusters_per_class=1, # number of cluster per class
          random_state=0)
```

```
[99]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
      ↪2,random_state=0)
```

6.5.1 Fitting decision tree classifier

```
[100]: model_dtc = DecisionTreeClassifier()
      model_dtc.fit(x_train,y_train)
      y_pred_dc = model_dtc.predict(x_test)
      print(f"Accuracy score of decision tree:
      ↪{round(accuracy_score(y_test,y_pred_dc)*100,2)}%")
```

Accuracy score of decision tree:77.5%

6.5.2 Bagging classifier

```
[101]: BC = BaggingClassifier(n_estimators = 100,random_state = 0)
      #n_estimator is number of trees we are fitting for bagging
      BC.fit(x_train,y_train)

      print(accuracy_score(BC.predict(x_test),y_test)*100)
```

85.0

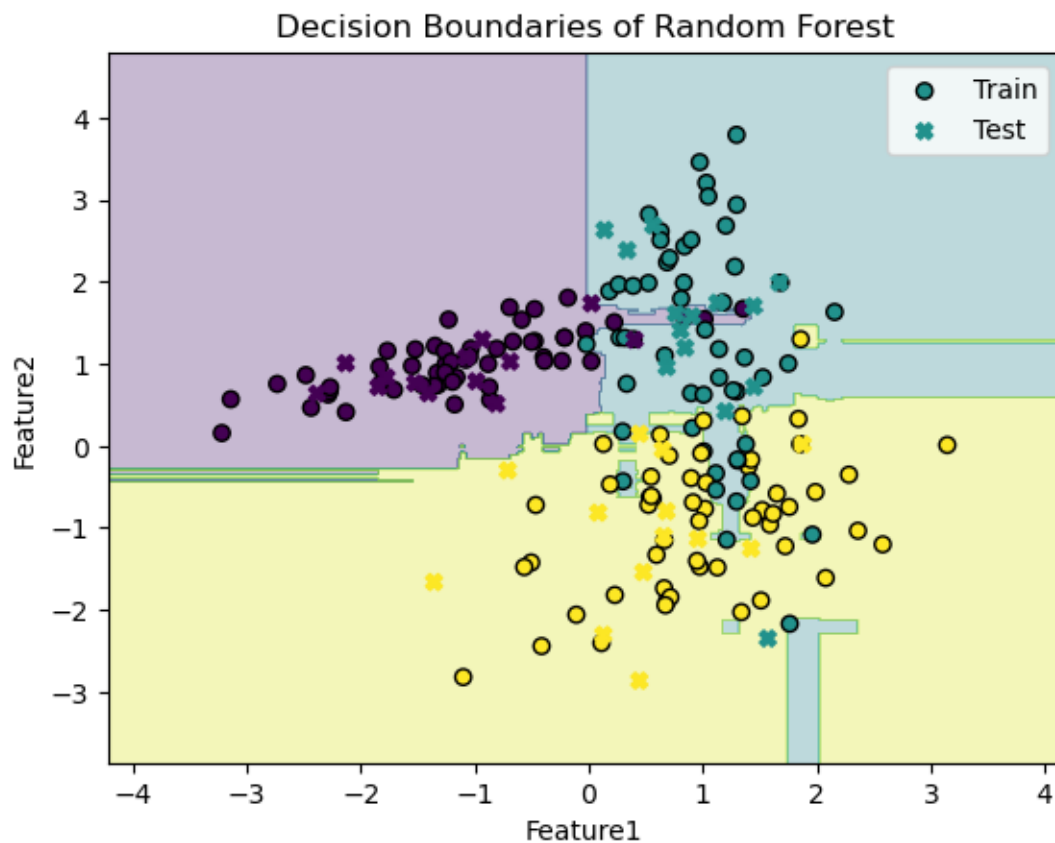
6.5.3 Random forest classifier

```
[103]: RFC = RandomForestClassifier(n_estimators=100,
      ↪max_features='sqrt',random_state=42)
      RFC.fit(x_train,y_train)
      print(accuracy_score(RFC.predict(x_test),y_test)*100)
```

87.5


6.6 Decision boundary for random forest

```
[104]: import matplotlib.pyplot as plt
x_min,x_max = x[:,0].min()-1, x[:,0].max()+1
y_min,y_max = x[:,1].min()-1,x[:,1].max()+1
xx,yy = np.meshgrid(np.linspace(x_min,x_max,500),np.linspace(y_min,y_max,500))
z = RFC.predict(np.c_[xx.ravel(),yy.ravel()])
z = z.reshape(xx.shape)
plt.contourf(xx,yy,z,alpha = 0.3)
plt.scatter(x_train[:,0],x_train[:,1],c = y_train,edgecolors='k',label = 'Train')
plt.scatter(x_test[:,0],x_test[:,1],c = y_test,marker='X',label = 'Test')
plt.xlabel("Feature1")
plt.ylabel("Feature2")
plt.title("Decision Boundaries of Random Forest")
plt.legend()
plt.show()
```



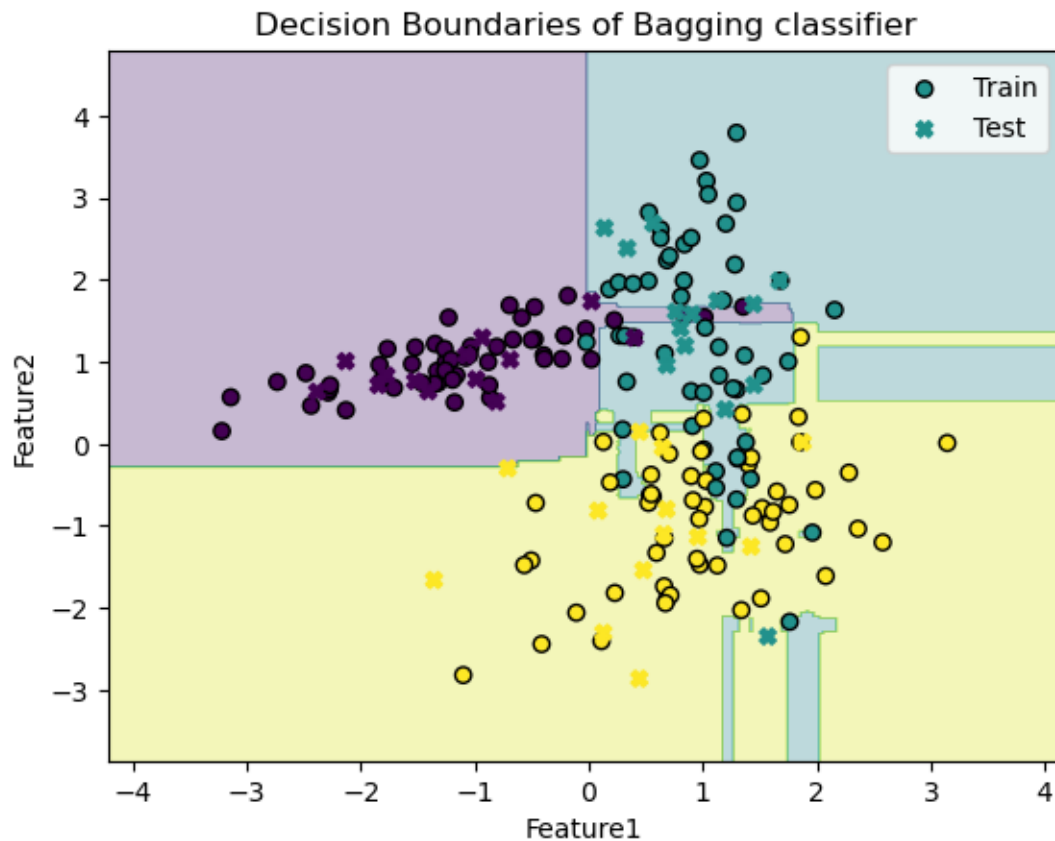
6.7 Decision boundary for Decision tree

```
[106]: x_min,x_max = x[:,0].min()-1, x[:,0].max()+1
y_min,y_max = x[:,1].min()-1,x[:,1].max()+1
xx,yy = np.meshgrid(np.linspace(x_min,x_max,500),np.linspace(y_min,y_max,500))
z = model_dtc.predict(np.c_[xx.ravel(),yy.ravel()])
z = z.reshape(xx.shape)
plt.contourf(xx,yy,z,alpha = 0.3)
plt.scatter(x_train[:,0],x_train[:,1],c = y_train,edgecolors='k',label = 'Train')
plt.scatter(x_test[:,0],x_test[:,1],c = y_test,marker='X',label = 'Test')
plt.xlabel("Feature1")
plt.ylabel("Feature2")
plt.title("Decision Boundaries of Decision Tree")
plt.legend()
plt.show()
```



6.8 Decision boundary for bagging classifier

```
[107]: x_min,x_max = x[:,0].min()-1, x[:,0].max()+1
y_min,y_max = x[:,1].min()-1,x[:,1].max()+1
xx,yy = np.meshgrid(np.linspace(x_min,x_max,500),np.linspace(y_min,y_max,500))
z = BC.predict(np.c_[xx.ravel(),yy.ravel()])
z = z.reshape(xx.shape)
plt.contourf(xx,yy,z,alpha = 0.3)
plt.scatter(x_train[:,0],x_train[:,1],c = y_train,edgecolors='k',label = 'Train')
plt.scatter(x_test[:,0],x_test[:,1],c = y_test,marker='X',label = 'Test')
plt.xlabel("Feature1")
plt.ylabel("Feature2")
plt.title("Decision Boundaries of Bagging classifier")
plt.legend()
plt.show()
```



6.9 Voting Classifier

```
[110]: from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression

log_clf = LogisticRegression()
rnd_clf = RandomForestClassifier(random_state=0)
sm_clf = LogisticRegression(multi_class='multinomial')
voting_clf = VotingClassifier(estimators=[('lr', log_clf), ('rf', rnd_clf), ('sm', sm_clf)],
                             voting = 'hard')
voting_clf.fit(x_train, y_train)
print(accuracy_score(y_test, voting_clf.predict(x_test))*100)
```

82.5

6.10 Change to voting criteria to soft by setting voting = 'soft' and check the output in the above case. In the case of soft voting the models output a probability that is averaged and used for prediction

```
[111]: from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression

log_clf = LogisticRegression()
rnd_clf = RandomForestClassifier(random_state=0)
sm_clf = LogisticRegression(multi_class='multinomial')
voting_clf = VotingClassifier(estimators=[('lr', log_clf), ('rf', rnd_clf), ('sm', sm_clf)],
                             voting = 'soft')
voting_clf.fit(x_train, y_train)
print(accuracy_score(y_test, voting_clf.predict(x_test))*100)
```

90.0

7 Adaboost Classifier

```
[117]: from sklearn.ensemble import AdaBoostClassifier
base_model = DecisionTreeClassifier(max_depth=1)
ABC = AdaBoostClassifier(estimator=base_model, n_estimators=500, random_state=0)
ABC.fit(x_train, y_train)
pred_ABC = ABC.predict(x_test)
print("AdaBoost Accuracy:", accuracy_score(y_test, pred_ABC)*100)
```

AdaBoost Accuracy: 87.5

7.1 Hyper parameter tuning for Randomforest

```
[124]: from sklearn.model_selection import GridSearchCV
param_grid = {
    'n_estimators': [10, 50, 100, 200, 300, 400, 500, 700],
    'max_depth': [3,5,7],
    'min_samples_leaf': [1, 2, 4]
}
model = RandomForestClassifier(random_state=0)
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5,
    ↪scoring='accuracy', n_jobs=-1)
grid_search.fit(x_train, y_train)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
y_pred = best_model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print("Best Parameters:", best_params)
print("Best Cross-validation Accuracy:", grid_search.best_score_)
print("Test Accuracy:", accuracy)
```

Best Parameters: {'max_depth': 5, 'min_samples_leaf': 4, 'n_estimators': 50}
Best Cross-validation Accuracy: 0.875
Test Accuracy: 0.95

7.2 Hyper parameter tuning for Adaboost

```
[128]: from sklearn.model_selection import GridSearchCV
param_grid = {
    'n_estimators': [10, 50, 75,100,125,150, 200, 400, 500],
    'estimator__max_depth':[1,2,3]
}
base_model = DecisionTreeClassifier(max_depth=1)
model = AdaBoostClassifier(estimator=base_model, random_state=0)
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5,
    ↪scoring='accuracy', n_jobs=-1)
grid_search.fit(x_train, y_train)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
y_pred = best_model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print("Best Parameters:", best_params)
print("Best Cross-validation Accuracy:", grid_search.best_score_)
print("Test Accuracy:", accuracy*100)
```

Best Parameters: {'estimator__max_depth': 3, 'n_estimators': 50}
Best Cross-validation Accuracy: 0.8625
Test Accuracy: 80.0

7.3 Now you can try to fit the multiple regression model, DecisionTreeRegressor, BaggingRegressor, RandomForestRegressor and AdaboostRegressor on the Book1.csv file and use the mean squared error to see how these ensemble models perform compared to the basic models.

```
[131]: ## Loading the dataset
df = pd.read_csv(r"C:\Users\Batch1\Documents\Downloads\TK\12_feb\Book1.csv")
df.drop('furnishingstatus',axis=1,inplace = True)
```

```
[132]: df
```

```
[132]:
```

	price	area	bedrooms	bathrooms	stories	parking
0	13300000	7420	4	2	3	2
1	12250000	8960	4	4	4	3
2	12250000	9960	3	2	2	2
3	12215000	7500	4	2	2	3
4	11410000	7420	4	1	2	2
..
244	4550000	5320	3	1	2	0
245	4550000	5360	3	1	2	2
246	4550000	3520	3	1	1	0
247	4550000	8400	4	1	4	3
248	4543000	4100	2	2	1	0

[249 rows x 6 columns]

```
[133]: scaler = MinMaxScaler()
X = scaler.fit_transform(df)
x = X[:,1:]
y = X[:,0]
```

```
[135]: ## Train test splitting
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
↪2,random_state=0)

## Fitting basic linear regression model
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
model_lin = LinearRegression()
model_lin.fit(x_train,y_train)
y_pred = model_lin.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.020077937566470735

r2 score: 0.18380110609849787

7.4 Using decision tree

```
[139]: from sklearn.tree import DecisionTreeRegressor
model_dtr = DecisionTreeRegressor()
model_dtr.fit(x_train,y_train)
y_pred = model_dtr.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.033563134393453355

r2 score: -0.36439278571874745

7.5 Using BaggingRegressor

```
[140]: from sklearn.ensemble import BaggingRegressor
model_bag = BaggingRegressor()
model_bag.fit(x_train,y_train)
y_pred = model_bag.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.023058824708278147

r2 score: 0.06262348115900129

7.6 Using pasting

```
[142]: from sklearn.ensemble import BaggingRegressor
model_paste = BaggingRegressor(bootstrap=False)
model_paste.fit(x_train,y_train)
y_pred = model_paste.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.03347673585441088

r2 score: -0.3608805528627723

7.7 Using randomforest regressor

```
[143]: from sklearn.ensemble import RandomForestRegressor
model_rf = RandomForestRegressor(n_estimators=100,
    ↪max_features='sqrt',random_state=0)
model_rf.fit(x_train,y_train)
y_pred = model_rf.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.02006166312349716

r2 score: 0.1844626871154682

7.8 Using Adaboost regressor

```
[144]: from sklearn.ensemble import AdaBoostRegressor
base_model = DecisionTreeRegressor(max_depth=1)
ABR = AdaBoostRegressor(estimator=base_model,n_estimators=500, random_state=0)
ABR.fit(x_train,y_train)
y_pred = ABR.predict(x_test)
print("MSE:",mean_squared_error(y_pred,y_test))
print("r2 score:",r2_score(y_test,y_pred))
```

MSE: 0.02265054829239708

r2 score: 0.0792205423832919