

Financial_Annalytics_LA3

April 5, 2025

Name	Tufan Kundu
Reg No.	24MDT0184
Course Name	Financial Analytics Lab
Course Code	PMDS610P
Assessment	Lab Digital Assessment 3

1 Problem Statement

You are given historical daily closing prices of the NIFTY 50 index from the National Stock Exchange of India (NSE). Your task is to analyze and forecast stock prices using an Regression model with time series error.

1.0.1 Importing the necessary libraries

```
[2]: import numpy as np
import pandas as pd
import yfinance as yf
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error
import warnings
warnings.filterwarnings('ignore')
```

```
[3]: # -----
# 2.1 Data Preprocessing
# -----
nifty = yf.download("^NSEI", start="2019-01-01", end="2024-01-01")
nifty = nifty[['Close']]
nifty.index = pd.to_datetime(nifty.index)
print("Missing values before handling:", nifty['Close'].isna().sum())
nifty.dropna(inplace=True)
print("Missing values after handling:", nifty['Close'].isna().sum())
```

YF.download() has changed argument auto_adjust default to True

[*****100%*****] 1 of 1 completed

Missing values before handling: Ticker

^NSEI 0

dtype: int64

Missing values after handling: Ticker

^NSEI 0

dtype: int64

```
[4]: # -----
# 2.2 Stationarity Check
# -----
def adf_test(series):
    result = adfuller(series)
    print(f"ADF Statistic: {result[0]:.4f}")
    print(f"p-value: {result[1]:.4f}")
    return result[1]

p_val = adf_test(nifty['Close'])

if p_val > 0.05:
    print("Series is non-stationary. Applying first-order differencing...")
    nifty_diff = nifty['Close'].diff().dropna()
    adf_test(nifty_diff)
else:
    print("Series is stationary.")
    nifty_diff = nifty['Close']
```

ADF Statistic: -0.1888

p-value: 0.9398

Series is non-stationary. Applying first-order differencing...

ADF Statistic: -12.2918

p-value: 0.0000

```
[5]: # -----
# 2.3 Model Selection & Training
# -----
nifty['Time'] = np.arange(len(nifty))
split = int(len(nifty) * 0.8)
train = nifty.iloc[:split]
test = nifty.iloc[split:]

X_train = train[['Time']]
y_train = train['Close']
X_test = test[['Time']]
y_test = test['Close']
```

```

lr = LinearRegression()
lr.fit(X_train, y_train)

y_pred_train = lr.predict(X_train)
residuals = y_train - y_pred_train

```

```

[6]: # Check stationarity of residuals
adf_residuals = adfuller(residuals)
print(f"\nADF Test on Residuals:")
print(f"ADF Statistic: {adf_residuals[0]}")
print(f"p-value: {adf_residuals[1]}")

# Fit ARMA model to residuals (ARIMA with d=0)
arma_model = ARIMA(residuals, order=(2, 0, 2)).fit()
print("\nARMA Model Summary:")
print(arma_model.summary())

```

ADF Test on Residuals:
ADF Statistic: -2.1465498132914145
p-value: 0.22621307051827794

ARMA Model Summary:

```

                                SARIMAX Results
=====
Dep. Variable:                  ^NSEI      No. Observations:                  985
Model:                         ARIMA(2, 0, 2)  Log Likelihood                  -6406.601
Date:                         Sat, 05 Apr 2025    AIC                           12825.201
Time:                         15:43:41          BIC                           12854.557
Sample:                        0                HQIC                       12836.367
                                - 985
Covariance Type:                opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	0.0011	688.632	1.66e-06	1.000	-1349.693	1349.696
ar.L1	0.1253	0.168	0.747	0.455	-0.204	0.454
ar.L2	0.8611	0.166	5.199	0.000	0.536	1.186
ma.L1	0.8673	0.171	5.082	0.000	0.533	1.202
ma.L2	-0.0175	0.025	-0.707	0.480	-0.066	0.031
sigma2	2.612e+04	693.795	37.653	0.000	2.48e+04	2.75e+04

```

=====
===
Ljung-Box (L1) (Q):                0.00    Jarque-Bera (JB):
1446.83
Prob(Q):                          0.98    Prob(JB):
0.00
Heteroskedasticity (H):            1.16    Skew:

```

-0.92
Prob(H) (two-sided): 0.18 Kurtosis:
8.65

=====
===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
[7]: # -----  
# 2.4 Model Evaluation  
# -----  
y_pred_reg = lr.predict(X_test).ravel()  
residuals_forecast = arma_model.forecast(steps=len(X_test))  
y_pred_final = y_pred_reg + residuals_forecast.values  
  
mse = mean_squared_error(y_test, y_pred_final)  
mape = mean_absolute_percentage_error(y_test, y_pred_final)  
  
print(f"\nModel Evaluation:")  
print(f"Mean Squared Error (MSE): {mse:.2f}")  
print(f"Mean Absolute Percentage Error (MAPE): {mape * 100:.2f}%")
```

Model Evaluation:

Mean Squared Error (MSE): 590709.10

Mean Absolute Percentage Error (MAPE): 3.45%

```
[8]: # -----  
# 2.5 Forecasting Future Prices  
# -----  
future_time = np.arange(len(nifty), len(nifty) + 30).reshape(-1, 1)  
future_reg = lr.predict(future_time).ravel()  
future_resid = arma_model.forecast(steps=30)  
forecast = future_reg + future_resid.values  
  
forecast_dates = pd.date_range(start=nifty.index[-1] + pd.Timedelta(days=1),  
                                ↪periods=30, freq='B')  
  
# -----  
# Display Forecasted Future Prices  
# -----  
forecast_df = pd.DataFrame({  
    'Date': forecast_dates,  
    'Forecasted_Close_Price': forecast  
)  
forecast_df.set_index('Date', inplace=True)
```

```
print("\nForecasted NIFTY 50 Closing Prices for the Next 30 Business Days:")
print(forecast_df.round(2))
```

Forecasted NIFTY 50 Closing Prices for the Next 30 Business Days:

Date	Forecasted_Close_Price
2024-01-01	20367.45
2024-01-02	20373.51
2024-01-03	20388.89
2024-01-04	20396.16
2024-01-05	20410.43
2024-01-08	20418.60
2024-01-09	20432.04
2024-01-10	20440.88
2024-01-11	20453.69
2024-01-12	20463.03
2024-01-15	20475.35
2024-01-16	20485.06
2024-01-17	20497.00
2024-01-18	20506.98
2024-01-19	20518.64
2024-01-22	20528.82
2024-01-23	20540.26
2024-01-24	20550.58
2024-01-25	20561.84
2024-01-26	20572.26
2024-01-29	20583.39
2024-01-30	20593.88
2024-01-31	20604.90
2024-02-01	20615.44
2024-02-02	20626.37
2024-02-05	20636.93
2024-02-06	20647.79
2024-02-07	20658.37
2024-02-08	20669.17
2024-02-09	20679.76

```
[9]: # -----
# Visualization (Focus on last few weeks)
# -----
last_n_days = 60
start_date = nifty.index[-last_n_days]

plt.figure(figsize=(12, 6))
```

```

plt.plot(nifty.index[-last_n_days:], nifty['Close'][-last_n_days:],  

        label='Historical Prices', color='blue')  

plt.plot(test.index[-len(y_pred_final):], y_pred_final, label='Test  

        Predictions', color='orange')  

plt.plot(forecast_dates, forecast, label='30-Day Forecast', color='green',  

        linestyle='--')  

plt.xlabel("Date")  

plt.ylabel("NIFTY 50 Closing Price")  

plt.title("NIFTY 50 Forecasting - Regression with Time Series Error (Last 12  

        Weeks)")  

plt.legend()  

plt.grid(True)  

plt.tight_layout()  

plt.xticks(rotation=45)  

plt.xlim(start_date, forecast_dates[-1])  

plt.show()

```

