

Financial Analytics Theory DA2

April 5, 2025

Name	Tufan Kundu
Reg No.	24MDT0184
Course Name	Financial Analytics
Course Code	PMDS610L
Assessment	Digital Assignment 2

1 Problem Statement

Predict the NIFTY 50 index values for the next 10 days using historical data from the past 3 years. The prediction should be performed by implementing and comparing the following time series models:

ARIMA (Autoregressive Integrated Moving Average), Regression with Time Series Error, ARFIMA (Autoregressive Fractionally Integrated Moving Average), GARCH (Generalized Autoregressive Conditional Heteroskedasticity)

1.0.1 Import necessary Libraries

```
[1]: import yfinance as yf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.ar_model import AutoReg
from arch import arch_model
from pmdarima import auto_arima
from pmdarima.arima.utils import ndiffs
from sklearn.metrics import mean_squared_error
from math import sqrt
import warnings
warnings.filterwarnings("ignore")
```

1.0.2 Load and prepare data

```
[2]: nifty_data = yf.download('^NSEI', period="3y", interval="1d")
df = nifty_data[['Close']].dropna()


# Create returns for volatility modeling
df['Returns'] = df['Close'].pct_change().dropna()

# Plot the data
plt.figure(figsize=(15, 8))
plt.plot(df['Close'], label='NIFTY 50 Closing Price')
plt.title('NIFTY 50 Closing Prices (Last 3 Years)')
plt.legend()
```


YF.download() has changed argument auto_adjust default to True

[*****100%*****] 1 of 1 completed

```
[2]: <matplotlib.legend.Legend at 0x180671a7a70>
```



```
[3]: plt.figure(figsize=(15, 8))
plt.plot(df['Returns'], label='Daily Returns', color='green')
plt.title('Daily Returns')
plt.legend()
plt.tight_layout()
plt.show()
```



1.0.3 Check stationarity using Adfuller

```
[4]: def test_stationarity(timeseries):
    result = adfuller(timeseries.dropna())
    print('ADF Statistic:', result[0])
    print('p-value:', result[1])
    print('Critical Values:', result[4])

    print("Stationarity test for Close prices:")
    test_stationarity(df['Close'])
```

Stationarity test for Close prices:

ADF Statistic: -0.7602317513297161

p-value: 0.8305255752841794

Critical Values: {'1%': -3.4392057325732104, '5%': -2.8654483492874236, '10%': -2.5688512291811225}

```
[5]: # Difference to make stationary
df['Close_diff'] = df['Close'].diff().dropna()
print("\nStationarity test for differenced series:")
test_stationarity(df['Close_diff'])


# Determine ARIMA parameters
plot_acf(df['Close_diff'].dropna(), lags=20)
plot_pacf(df['Close_diff'].dropna(), lags=20)
plt.show()
```


Stationarity test for differenced series:

ADF Statistic: -28.17810031286188

p-value: 0.0

Critical Values: {'1%': -3.4392177416762086, '5%': -2.8654536415414684, '10%': -2.5688540485756026}





1.1 Arima

1.1.1 Auto ARIMA to find best parameters

```
[6]: auto_model = auto_arima(df['Close'], seasonal=False,
    ↪ trace=True, error_action='ignore', suppress_warnings=True)
    print(auto_model.summary())
```

Performing stepwise search to minimize aic

```
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=9711.973, Time=1.63 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=9707.644, Time=0.03 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=9708.379, Time=0.05 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=9708.503, Time=0.17 sec
ARIMA(0,1,0)(0,0,0)[0]           : AIC=9706.758, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=9708.250, Time=1.23 sec
```

Best model: ARIMA(0,1,0)(0,0,0)[0]

Total fit time: 3.151 seconds

SARIMAX Results

```
=====
Dep. Variable:          y    No. Observations:          742
Model:                SARIMAX(0, 1, 0)    Log Likelihood    -4852.379
```

Date: Sat, 05 Apr 2025 AIC 9706.758
Time: 22:03:09 BIC 9711.366
Sample: 0 HQIC 9708.535
- 742

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2      2.85e+04    700.216    40.700    0.000    2.71e+04    2.99e+04
=====
```

===

Ljung-Box (L1) (Q): 1.13 Jarque-Bera (JB): 1591.95
Prob(Q): 0.29 Prob(JB): 0.00
Heteroskedasticity (H): 1.83 Skew: -0.69
Prob(H) (two-sided): 0.00 Kurtosis: 10.05

```
=====
===
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

1.1.2 Splitting the Data

```
[7]: # Split train-test (last 30 days for testing)
train = df['Close'][:-30]
test = df['Close'][-30:]

# Fit ARIMA model (using parameters from auto_arima)
model_arima = ARIMA(train, order=(2,1,2))
model_arima_fit = model_arima.fit()
print(model_arima_fit.summary())

# Forecast
forecast = model_arima_fit.forecast(steps=30)
rmse = sqrt(mean_squared_error(test, forecast))
print(f'ARIMA RMSE: {rmse}')
```

SARIMAX Results

```
=====
Dep. Variable:          ^NSEI    No. Observations:          712
Model:                ARIMA(2, 1, 2)    Log Likelihood          -4650.678
Date:                Sat, 05 Apr 2025    AIC                      9311.355
Time:                22:03:10    BIC                      9334.188
```


Sample: 0 HQIC 9320.175
 - 712
 Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.7426	4.027	-0.184	0.854	-8.635	7.149
ar.L2	-0.0204	3.079	-0.007	0.995	-6.055	6.014
ma.L1	0.6964	4.018	0.173	0.862	-7.178	8.571
ma.L2	0.0304	2.870	0.011	0.992	-5.594	5.655
sigma2	2.811e+04	922.877	30.460	0.000	2.63e+04	2.99e+04

=====
 ===
 Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB):
 1490.45
 Prob(Q): 1.00 Prob(JB):
 0.00
 Heteroskedasticity (H): 1.74 Skew:
 -0.74
 Prob(H) (two-sided): 0.00 Kurtosis:
 9.94
 =====
 ===

Warnings:
 [1] Covariance matrix calculated using the outer product of gradients (complex-
 step).
 ARIMA RMSE: 484.6059918474232

```
[8]: # Plot
plt.figure(figsize=(12,6))
plt.plot(train.index[-100:], train[-100:], label='Train')
plt.plot(test.index, test, label='Actual')
plt.plot(test.index, forecast, label='ARIMA Forecast', color='red')
plt.title('ARIMA Forecast vs Actual')
plt.legend()
plt.show()
```




1.2 Regression with Time Series

1.2.1 Determine optimal lags using PACF

```
[9]: plot_pacf(df['Close'], lags=20)
plt.show()

# Fit AutoReg model (using 5 lags based on PACF)
model_autoreg = AutoReg(train, lags=5)
model_autoreg_fit = model_autoreg.fit()
print(model_autoreg_fit.summary())

# Forecast
forecast_autoreg = model_autoreg_fit.predict(start=len(train),
end=len(train)+29)
rmse_autoreg = sqrt(mean_squared_error(test, forecast_autoreg))
print(f'AutoReg RMSE: {rmse_autoreg}')
```

AutoReg Model Results

```

=====
Dep. Variable:          ^NSEI    No. Observations:          712
Model:                  AutoReg(5)  Log Likelihood             -4623.779
Method:                  Conditional MLE  S.D. of innovations         167.513
Date:                    Sat, 05 Apr 2025  AIC                        9261.558
Time:                    22:03:10    BIC                        9293.485
Sample:                  5    HQIC                        9273.894
                        712
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	41.7769	44.747	0.934	0.350	-45.925	129.479
^NSEI.L1	0.9494	0.038	25.284	0.000	0.876	1.023
^NSEI.L2	0.0945	0.052	1.827	0.068	-0.007	0.196
^NSEI.L3	-0.0406	0.052	-0.784	0.433	-0.142	0.061
^NSEI.L4	0.0434	0.052	0.840	0.401	-0.058	0.145
^NSEI.L5	-0.0484	0.038	-1.291	0.197	-0.122	0.025

Roots

```


=====
Real          Imaginary      Modulus      Frequency
=====

```

AR.1	1.0018	-0.0000j	1.0018	-0.0000
AR.2	-1.9906	-0.0000j	1.9906	-0.5000
AR.3	2.0542	-0.0000j	2.0542	-0.0000
AR.4	-0.0840	-2.2442j	2.2458	-0.2560
AR.5	-0.0840	+2.2442j	2.2458	0.2560

AutoReg RMSE: 479.3048230892029

```
[10]: # Plot
plt.figure(figsize=(12,6))
plt.plot(train.index[-100:], train[-100:], label='Train')
plt.plot(test.index, test, label='Actual')
plt.plot(test.index, forecast_autoreg, label='AutoReg Forecast', color='green')
plt.title('AutoReg Forecast vs Actual')
plt.legend()
plt.show()
```



1.3 Automated ARFIMA modeling

```
[11]: close_series = nifty_data['Close'].dropna()
print("Searching for best ARFIMA parameters...")
model = auto_arima(close_series,
                    start_p=0, max_p=3,      # AR order
                    start_q=0, max_q=3,      # MA order
                    d=None,                  # Let model determine differencing
                    test='adf',              # Augmented Dickey-Fuller test
                    seasonal=False,          # No seasonality
```

```

        trace=True,                # Print progress
        error_action='ignore',
        suppress_warnings=True,
        stepwise=True)             # Smart parameter search

print("\nBest model summary:")
print(model.summary())

```

Searching for best ARFIMA parameters...

Performing stepwise search to minimize aic

```

ARIMA(0,1,0)(0,0,0)[0] intercept    : AIC=9707.644, Time=0.04 sec
ARIMA(1,1,0)(0,0,0)[0] intercept    : AIC=9708.379, Time=0.10 sec
ARIMA(0,1,1)(0,0,0)[0] intercept    : AIC=9708.503, Time=0.31 sec
ARIMA(0,1,0)(0,0,0)[0]               : AIC=9706.758, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept    : AIC=9708.250, Time=1.46 sec

```

Best model: ARIMA(0,1,0)(0,0,0)[0]

Total fit time: 1.954 seconds

Best model summary:

SARIMAX Results

```

=====
Dep. Variable:          y      No. Observations:          742
Model:                SARIMAX(0, 1, 0)  Log Likelihood      -4852.379
Date:                 Sat, 05 Apr 2025   AIC                  9706.758
Time:                 22:03:13          BIC                  9711.366
Sample:               0                HQIC                  9708.535
                   - 742

```

Covariance Type: opg

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2      2.85e+04    700.216    40.700      0.000    2.71e+04    2.99e+04
=====

```

```

===
Ljung-Box (L1) (Q):          1.13   Jarque-Bera (JB):
1591.95
Prob(Q):                     0.29   Prob(JB):
0.00
Heteroskedasticity (H):      1.83   Skew:
-0.69
Prob(H) (two-sided):         0.00   Kurtosis:
10.05
=====

```


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-

step).

```
[12]: # Forecast next 30 days
forecast, conf_int = model.predict(n_periods=30, return_conf_int=True)
forecast_index = pd.date_range(start=close_series.index[-1] + pd.
    ↪Timedelta(days=1), periods=30)
```

```
[13]: # Plot results
plt.figure(figsize=(12,6))
plt.plot(close_series.index[-100:], close_series.values[-100:],
    ↪label='Observed')
plt.plot(forecast_index, forecast, label='Forecast', color='red')
plt.fill_between(forecast_index,
    conf_int[:,0],
    conf_int[:,1],
    color='pink', alpha=0.3, label='95% CI')
plt.title('ARFIMA Forecast for NIFTY 50')
plt.xlabel('Date')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()
```



```
[14]: # Print forecasted values
forecast_df = pd.DataFrame({
    'Date': forecast_index,
    'Forecast': forecast,
```

```

    'Lower CI': conf_int[:,0],
    'Upper CI': conf_int[:,1]
})
print("\n30-day Forecast:")
print(forecast_df.round(2))

```

30-day Forecast:

	Date	Forecast	Lower CI	Upper CI
742	2025-04-05	22904.45	22573.58	23235.32
743	2025-04-06	22904.45	22436.53	23372.37
744	2025-04-07	22904.45	22331.36	23477.54
745	2025-04-08	22904.45	22242.70	23566.19
746	2025-04-09	22904.45	22164.60	23644.30
747	2025-04-10	22904.45	22093.98	23714.92
748	2025-04-11	22904.45	22029.04	23779.86
749	2025-04-12	22904.45	21968.60	23840.30
750	2025-04-13	22904.45	21911.83	23897.07
751	2025-04-14	22904.45	21858.14	23950.76
752	2025-04-15	22904.45	21807.07	24001.83
753	2025-04-16	22904.45	21758.27	24050.62
754	2025-04-17	22904.45	21711.47	24097.43
755	2025-04-18	22904.45	21666.44	24142.46
756	2025-04-19	22904.45	21622.99	24185.91
757	2025-04-20	22904.45	21580.96	24227.94
758	2025-04-21	22904.45	21540.23	24268.67
759	2025-04-22	22904.45	21500.68	24308.22
760	2025-04-23	22904.45	21462.21	24346.69
761	2025-04-24	22904.45	21424.74	24384.16
762	2025-04-25	22904.45	21388.20	24420.70
763	2025-04-26	22904.45	21352.52	24456.38
764	2025-04-27	22904.45	21317.64	24491.26
765	2025-04-28	22904.45	21283.51	24525.39
766	2025-04-29	22904.45	21250.09	24558.81
767	2025-04-30	22904.45	21217.32	24591.57
768	2025-05-01	22904.45	21185.19	24623.71
769	2025-05-02	22904.45	21153.64	24655.26
770	2025-05-03	22904.45	21122.65	24686.25
771	2025-05-04	22904.45	21092.19	24716.71

1.3.1 GARCH

```
[15]: returns = df['Returns'].dropna()
```

1.3.2 Splitting the Data

```
[16]: # Split train-test
train_ret = returns[:-30]
test_ret = returns[-30:]

# Fit GARCH(1,1) model
model_garch = arch_model(train_ret, mean='Zero', vol='GARCH', p=1, q=1)
model_garch_fit = model_garch.fit(update_freq=5)
print(model_garch_fit.summary())

# Forecast volatility
forecast_garch = model_garch_fit.forecast(horizon=30)
forecast_volatility = np.sqrt(forecast_garch.variance.values[-1,:])


Optimization terminated successfully      (Exit mode 0)
      Current function value: -2436.5008155105215
      Iterations: 7
      Function evaluations: 17
      Gradient evaluations: 3

      Zero Mean - GARCH Model Results
=====
Dep. Variable:          Returns      R-squared:                0.000
Mean Model:            Zero Mean    Adj. R-squared:           0.001
Vol Model:             GARCH        Log-Likelihood:          2436.50
Distribution:          Normal       AIC:                    -4867.00
Method:               Maximum Likelihood BIC:                    -4853.30
                                     No. Observations:          711
Date:                 Sat, Apr 05 2025 Df Residuals:              711
Time:                 22:03:13      Df Model:                  0
                                     Volatility Model
=====
               coef      std err          t      P>|t|      95.0% Conf. Int.
-----
omega      1.9519e-05  2.950e-12  6.618e+06      0.000 [1.952e-05,1.952e-05]
alpha[1]      0.1957  6.633e-02    2.950  3.177e-03  [6.568e-02, 0.326]
beta[1]      0.5179  4.917e-02   10.534  6.036e-26  [ 0.422, 0.614]
=====

Covariance estimator: robust
```

```
[17]: # Plot
plt.figure(figsize=(12,6))
plt.plot(test_ret.index, test_ret, label='Actual Returns')
plt.plot(test_ret.index, forecast_volatility, label='Forecast Volatility',
         color='red')
plt.plot(test_ret.index, -forecast_volatility, color='red')
plt.title('GARCH Volatility Forecast')
```

```
plt.legend()
plt.show()
```



1.4 Arima Forecast


1.4.1 Generate future dates (10 business days)

```
[18]: future_dates = pd.date_range(start=df.index[-1] + pd.Timedelta(days=1),
    ↪ periods=10, freq='B')
```

```
final_model = ARIMA(df['Close'], order=(2,1,2))
final_model_fit = final_model.fit()
future_forecast = final_model_fit.forecast(steps=10)
```

```
[19]: # Plot final forecast
plt.figure(figsize=(12,6))
plt.plot(df['Close'][-100:], label='Historical Prices')
plt.plot(future_dates, future_forecast, label='10-Day ARIMA Forecast',
    ↪ color='red', marker='o')
plt.title('NIFTY 50 - 10-Day Price Forecast')
plt.legend()
plt.grid(True)
plt.show()

print("Future Date Forecasts:")
print(pd.DataFrame({'Date': future_dates, 'Forecast': future_forecast}))
```



Future Date Forecasts:

	Date	Forecast
742	2025-04-07	22903.249786
743	2025-04-08	22893.288035
744	2025-04-09	22894.733804
745	2025-04-10	22890.672337
746	2025-04-11	22891.997009
747	2025-04-14	22890.200780
748	2025-04-15	22891.047034
749	2025-04-16	22890.202924
750	2025-04-17	22890.685631
751	2025-04-18	22890.272726