

# 7\_FEB\_Regression\_analysis\_hw

February 9, 2025

## 1 Experiment IV

## 2 Tufan Kundu

## 3 Reg no: 24MDT0184

### 3.1 Importing the libraries

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
```

### 3.2 Reading the dataset

```
[7]: df = pd.read_excel(r"D:\study material\VIT_Data_Science\Winter_Sem\Regression_
↳Analysis and Predictive Models Lab\7_feb\advertising.xlsx")
df
```

```
[7]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
..	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

[200 rows x 4 columns]

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   TV           200 non-null    float64
1   Radio        200 non-null    float64
2   Newspaper    200 non-null    float64
3   Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

### 3.3 Perform Linear Regression using sklearn:

- Define TV as the independent variable (X) and Sales as the dependent variable (y).

```
[10]: x_tv = df['TV'].values.reshape(-1,1)
      y = df['Sales'].values

      model_tv = LinearRegression()
      model_tv.fit(x_tv,y)
      y_pred_tv = model_tv.predict(x_tv)
      residuals_tv = y - y_pred_tv
```


#### 3.3.1 Scatter plot of TV vs Sales along with the regression line and the residuals

```
[12]: ## creating plot to visualize the regression line

plt.figure(figsize = (10,6))
plt.scatter(x = x_tv, y = y, color = 'blue', label = 'Data points', alpha = 0.5)
plt.plot(x_tv,y_pred_tv, color = 'red', label = "Regression Line")

## Adding lines showing the residuals(the vertical distance between actual and
↪predicted value)
for i in range(len(x_tv)):
    plt.plot([x_tv[i],x_tv[i]], [y[i],y_pred_tv[i]], color = 'grey', lw=1)

plt.title("Least square fit for sales vs TV advertising")
plt.xlabel("TV advertising budget(in thousands of dollars)")
plt.ylabel("Sales (in thousand of units)")
plt.legend()
plt.show()
```



### 3.4 Perform Linear Regression using statsmodels:

```
[14]: import statsmodels.api as sm
x_tv = df['TV']
y = df['Sales']

# adds a constant to the independent variable
x_with_const = sm.add_constant(x_tv) # adds a column of ones to x for the
↪ intercept form

# Fit the OLS regression model using stats model
model_sm_tv = sm.OLS(y,x_with_const).fit()

print(model_sm_tv.summary())
# get the 95% confidence interval for the model coefficients ( and )

confidence_intervals_tv = model_sm_tv.conf_int(alpha = 0.05)

# print the confidence intervals for the intercept and coefficient

print("95% confidence interval for :\n",confidence_intervals_tv.iloc[0])
print("95% confidence interval for 1:\n",confidence_intervals_tv.iloc[1])
```

```

standard_error_tv = model_sm_tv.bse

print("Standard error for (Intercept):",standard_error_tv.iloc[0])
print("Standard error for 1:",standard_error_tv.iloc[1])

```

#### OLS Regression Results

```

=====
Dep. Variable:          Sales    R-squared:                0.812
Model:                  OLS      Adj. R-squared:            0.811
Method:                 Least Squares    F-statistic:          856.2
Date:                  Sun, 09 Feb 2025    Prob (F-statistic):      7.93e-74
Time:                  11:02:44    Log-Likelihood:         -448.99
No. Observations:      200    AIC:                   902.0
Df Residuals:          198    BIC:                   908.6
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	6.9748	0.323	21.624	0.000	6.339	7.611
TV	0.0555	0.002	29.260	0.000	0.052	0.059

```

=====
Omnibus:                0.013    Durbin-Watson:          2.029
Prob(Omnibus):          0.993    Jarque-Bera (JB):        0.043
Skew:                  -0.018    Prob(JB):                0.979
Kurtosis:               2.938    Cond. No.:               338.
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

95% confidence interval for :

0 6.338740

1 7.610903

Name: const, dtype: float64

95% confidence interval for 1:

0 0.051727

1 0.059203

Name: TV, dtype: float64

Standard error for (Intercept): 0.3225534848524013

Standard error for 1: 0.001895551178040241

### 3.5 Simple Linear Regression for Radio advertising

```

[16]: x_radio = df['Radio'].values.reshape(-1,1)
      y = df['Sales'].values

```

```

model_radio = LinearRegression()
model_radio.fit(x_radio,y)
y_pred_radio = model_radio.predict(x_radio)
residuals_radio = y - y_pred_radio

```

### 3.5.1 Scatter plot of Radio vs Sales along with the regression line and the residuals

```


[18]: ## creating plot to visualize the regression line

plt.figure(figsize = (10,6))
plt.scatter(x = x_radio, y = y, color = 'blue', label = 'Data points', alpha = 0.5)
plt.plot(x_radio,y_pred_radio, color = 'red', label = "Regression Line")

## Adding lines showing the residuals(the vertical distance between actual and predicted value)
for i in range(len(x_radio)):
    plt.plot([x_radio[i],x_radio[i]], [y[i],y_pred_radio[i]], color = 'grey', lw=1)

plt.title("Least square fit for sales vs Radio advertising")
plt.xlabel("Radio advertising budget(in thousands of dollars)")
plt.ylabel("Sales (in thousand of units)")
plt.legend()
plt.show()

```



### 3.5.2 Performing Linear regression using stats model

```
[20]: import statsmodels.api as sm
x_radio = df['Radio']
y = df['Sales']

# adds a constant to the independent variable
x_with_const = sm.add_constant(x_radio) # adds a column of ones to x for the
↳ intercept form

# Fit the OLS regression model using stats model
model_sm_radio = sm.OLS(y,x_with_const).fit()

print(model_sm_radio.summary())

# get the 95% confidence interval for the model coefficients ( and )

confidence_intervals_radio = model_sm_radio.conf_int(alpha = 0.05)

# print the confidence intervals for the intercept and coefficient

print("95% confidence interval for :\n",confidence_intervals_radio.iloc[0])
print("95% confidence interval for 1:\n",confidence_intervals_radio.iloc[1])

standard_error_radio = model_sm_radio.bse

print("Standard error for (Intercept):",standard_error_radio.iloc[0])
print("Standard error for 1:",standard_error_radio.iloc[1])
```

#### OLS Regression Results

=====						
Dep. Variable:	Sales	R-squared:	0.122			
Model:	OLS	Adj. R-squared:	0.118			
Method:	Least Squares	F-statistic:	27.57			
Date:	Sun, 09 Feb 2025	Prob (F-statistic):	3.88e-07			
Time:	11:02:47	Log-Likelihood:	-603.18			
No. Observations:	200	AIC:	1210.			
Df Residuals:	198	BIC:	1217.			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	12.2357	0.653	18.724	0.000	10.947	13.524
Radio	0.1244	0.024	5.251	0.000	0.078	0.171

```
=====
Omnibus:                11.077    Durbin-Watson:                2.018
Prob(Omnibus):           0.004    Jarque-Bera (JB):         9.124
Skew:                   -0.433    Prob(JB):                 0.0104
Kurtosis:                2.414    Cond. No.                 51.4
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

95% confidence interval for :

0 10.947036

1 13.524408

Name: const, dtype: float64

95% confidence interval for 1:

0 0.077703

1 0.171161

Name: Radio, dtype: float64

Standard error for (Intercept): 0.653486294381421

Standard error for 1: 0.023696033393257657

### 3.6 Simple Linear Regression for Newspaper advertising

```
[22]: x_np = df['Newspaper'].values.reshape(-1,1)
      y = df['Sales'].values

      model_np = LinearRegression()
      model_np.fit(x_np,y)
      y_pred_np = model_np.predict(x_np)
      residuals_np = y - y_pred_np
```

#### 3.6.1 Scatter plot of Newspaper vs Sales along with the regression line and the residuals


```
[24]: ## creating plot to visualize the regression line

plt.figure(figsize = (10,6))
plt.scatter(x = x_np, y = y, color = 'blue', label = 'Data points', alpha = 0.5)
plt.plot(x_np,y_pred_np, color = 'red', label = "Regression Line")

## Adding lines showing the residuals(the vertical distance between actual and
↪predicted value)
for i in range(len(x_np)):
    plt.plot([x_np[i],x_np[i]], [y[i],y_pred_np[i]], color = 'grey', lw=1)

plt.title("Least square fit for sales vs Newspaper advertising")
plt.xlabel("Newspaper advertising budget(in thousands of dollars)")
```

```
plt.ylabel("Sales (in thousand of units)")
plt.legend()
plt.show()
```



### 3.6.2 Performing Linear regression using stats model

```
[26]: import statsmodels.api as sm
x_np = df['Newspaper']
y = df['Sales']

# adds a constant to the independent variable
x_with_const = sm.add_constant(x_np) # adds a column of ones to x for the
↪ intercept form

# Fit the OLS regression model using stats model
model_sm_np = sm.OLS(y,x_with_const).fit()

print(model_sm_np.summary())

# get the 95% confidence interval for the model coefficients ( and )

confidence_intervals_np = model_sm_np.conf_int(alpha = 0.05)
```



```
# print the confidence intervals for the intercept and coefficient

print("95% confidence interval for :\\n",confidence_intervals_np.iloc[0])
print("95% confidence interval for 1:\\n",confidence_intervals_np.iloc[1])

standard_error_np = model_sm_np.bse

print("Standard error for (Intercept):",standard_error_np.iloc[0])
print("Standard error for 1:",standard_error_np.iloc[1])
```

#### OLS Regression Results

```
=====
Dep. Variable:          Sales    R-squared:                0.025
Model:                  OLS      Adj. R-squared:            0.020
Method:                 Least Squares    F-statistic:          5.067
Date:                  Sun, 09 Feb 2025    Prob (F-statistic):    0.0255
Time:                  11:02:48    Log-Likelihood:       -613.69
No. Observations:      200    AIC:                  1231.
Df Residuals:          198    BIC:                  1238.
Df Model:               1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	13.9595	0.638	21.870	0.000	12.701	15.218
Newspaper	0.0383	0.017	2.251	0.025	0.005	0.072

```
=====
Omnibus:                 10.252    Durbin-Watson:          2.017
Prob(Omnibus):           0.006    Jarque-Bera (JB):        4.808
Skew:                    -0.111    Prob(JB):                0.0903
Kurtosis:                 2.273    Cond. No.                 64.7
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

95% confidence interval for :

0 12.700833

1 15.218265

Name: const, dtype: float64

95% confidence interval for 1:

0 0.004749

1 0.071899

Name: Newspaper, dtype: float64

Standard error for (Intercept): 0.6382885131624155

Standard error for 1: 0.017025666500722073

### 3.6.3 Comparison of the three models

```
[28]: print(" 0(Intercept) for TV:",model_tv.intercept_)
      print(" 1(Coefficient) for TV:",model_tv.coef_[0])
      print()
      print(" 0(Intercept) for Radio:",model_radio.intercept_)
      print(" 1(Coefficient) for Radio:",model_radio.coef_[0])
      print()
      print(" 0(Intercept) for Newspaper:",model_np.intercept_)
      print(" 1(Coefficient) for Newspaper:",model_np.coef_[0])
```

```
0(Intercept) for TV: 6.974821488229891
1(Coefficient) for TV: 0.055464770469558874
```

```
0(Intercept) for Radio: 12.235721966369233
1(Coefficient) for Radio: 0.12443165550338577
```

```
0(Intercept) for Newspaper: 13.959548653554414
1(Coefficient) for Newspaper: 0.03832399510524274
```

```
[29]: from sklearn.metrics import r2_score
      from sklearn.metrics import mean_squared_error
      print("R2 score for TV model:",r2_score(y,y_pred_tv))
      print("MSE for TV model:",mean_squared_error(y,y_pred_tv))
      print()
      print("R2 score for Radio model:",r2_score(y,y_pred_radio))
      print("MSE for Radio model:",mean_squared_error(y,y_pred_radio))
      print()
      print("R2 score for Newspaper model:",r2_score(y,y_pred_np))
      print("MSE for Newspaper model:",mean_squared_error(y,y_pred_np))
```

```
R2 score for TV model: 0.8121757029987414
MSE for TV model: 5.2177438977951285
```

```
R2 score for Radio model: 0.1222419039947863
MSE for Radio model: 24.384049466937633
```

```
R2 score for Newspaper model: 0.024951369862864836
MSE for Newspaper model: 27.086772697557045
```

### 3.7 Summarized Insights:

- **TV advertising has the highest impact on sales, explaining 81.2% of the variation in sales with the lowest prediction error (MSE = 5.22).**
- **Radio advertising has a weak impact on sales, explaining only 12.2% of the variation, with higher prediction errors (MSE = 24.38).**
- **Newspaper advertising has minimal influence on sales, explaining just 2.5% of the**

**variation**, with the highest prediction error ( $\text{MSE} = 27.09$ ).

- TV should be the primary focus for advertising investments as it provides the best return.
- Radio may be considered as a supplementary channel.
- Newspaper advertising is not a cost-effective medium for increasing sales and should be re-considered.